Supervised Learning

DATA SCIENCE BOOTCAMP

What is Supervised Learning?

Data with correct answers → model

New Data without answers → model → Predicted answers

Regression: "answers" are numeric

Movie data including gross → model

Movie data without gross → model → Predicted gross

Classification: "answers" are categories

Movie data including Oscar wins or lack thereof

model

Movie data without Oscar results

model

Predict winning oscar

Classification: "answers" are categories

Breast cancer surgery patient data including (survived/not)

model

Patient data without survival result

model

Predict survival

Classification: "answers" are categories

Color, shape, weight, sweetness, sourness for a bunch of apples, bananas & peaches

model

Color, shape, weight, sweetness, sourness (without fruit type)

model

Predict apple, banana or peach

Classification: “answers” are categories

Example: each data point (row)

Label: each category to be predicted

Feature: each property (column) used in predicting

1 Feature: Number of + nodes
2 Labels: Survived / Lost

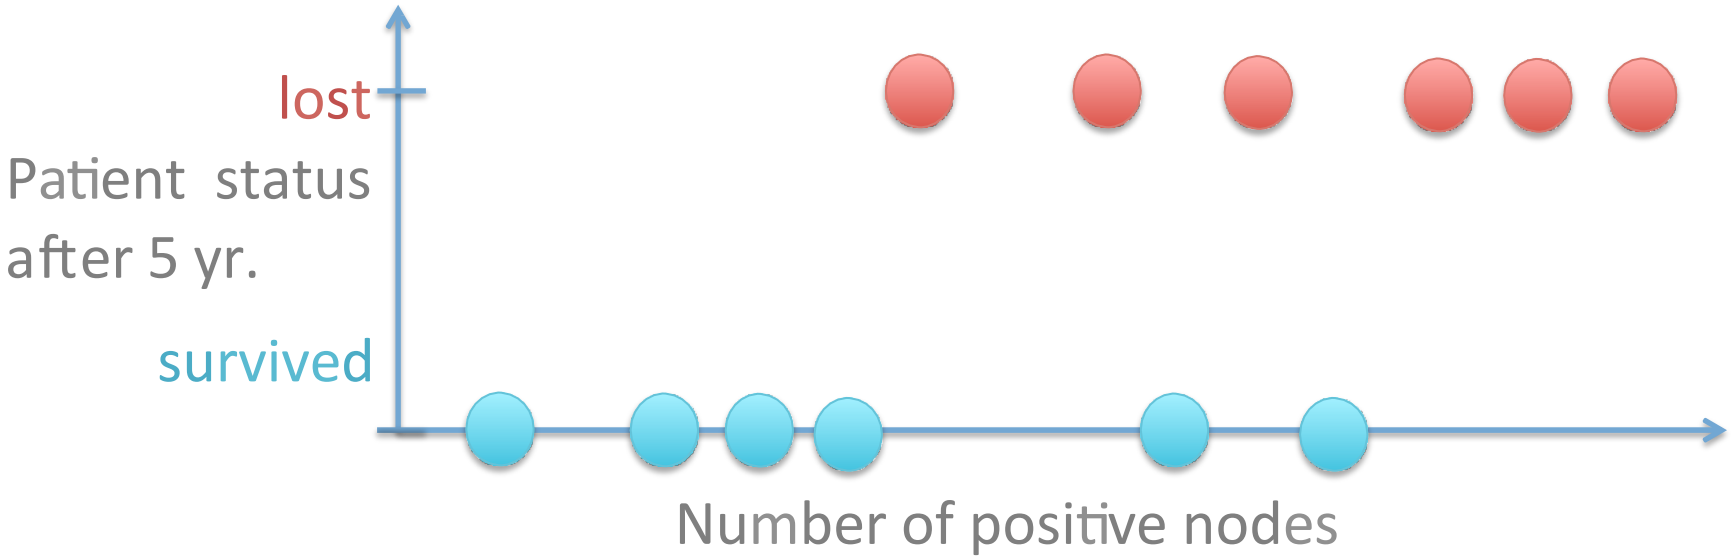

1 Feature: Number of + nodes

2 Labels: Survived / Lost

Number of positive nodes

1 Feature: Number of + nodes

2 Labels: Survived / Lost

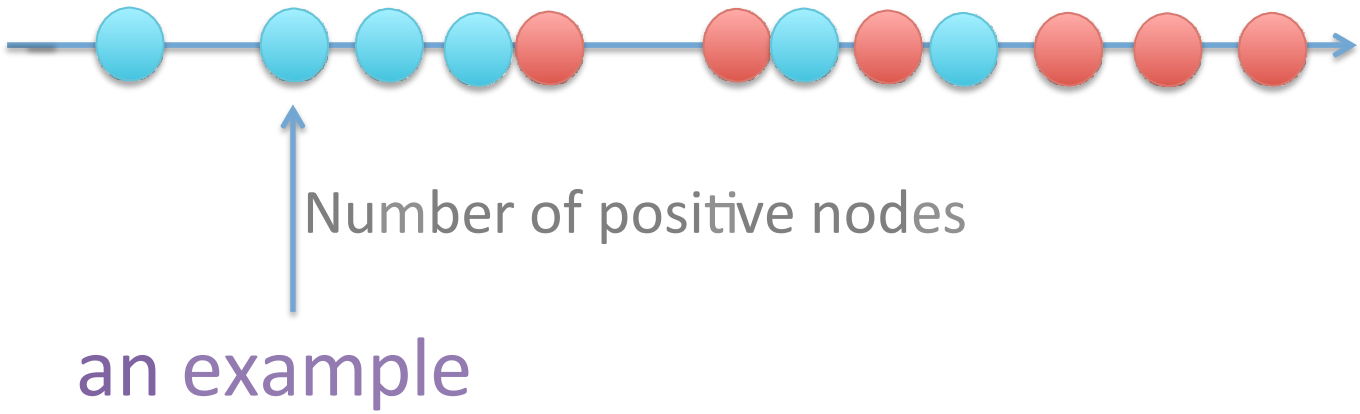

1 Feature: Number of + nodes

2 Labels: Survived / Lost

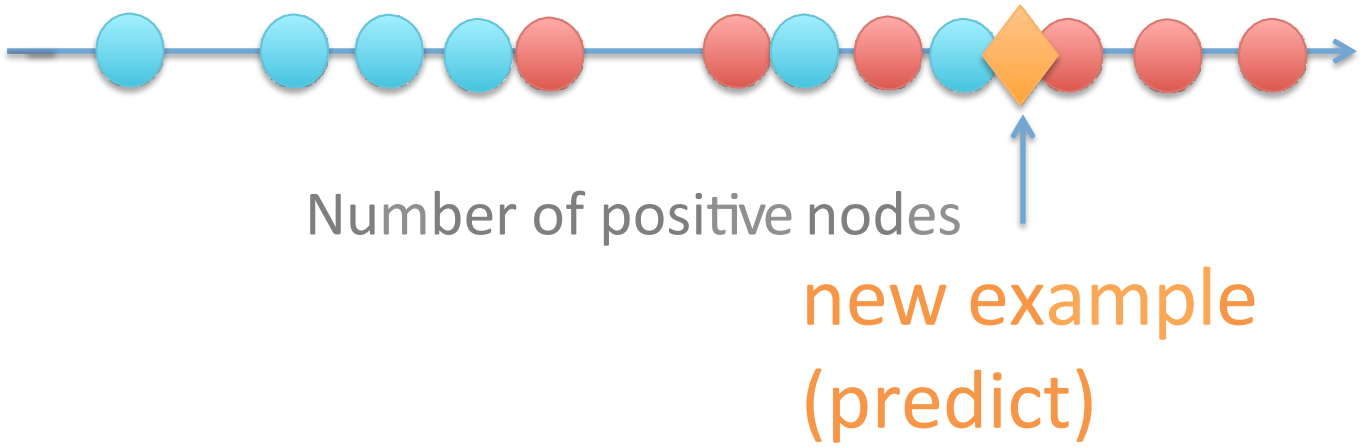

2 Features: Number of + nodes, Age

2 Labels: Survived / Lost

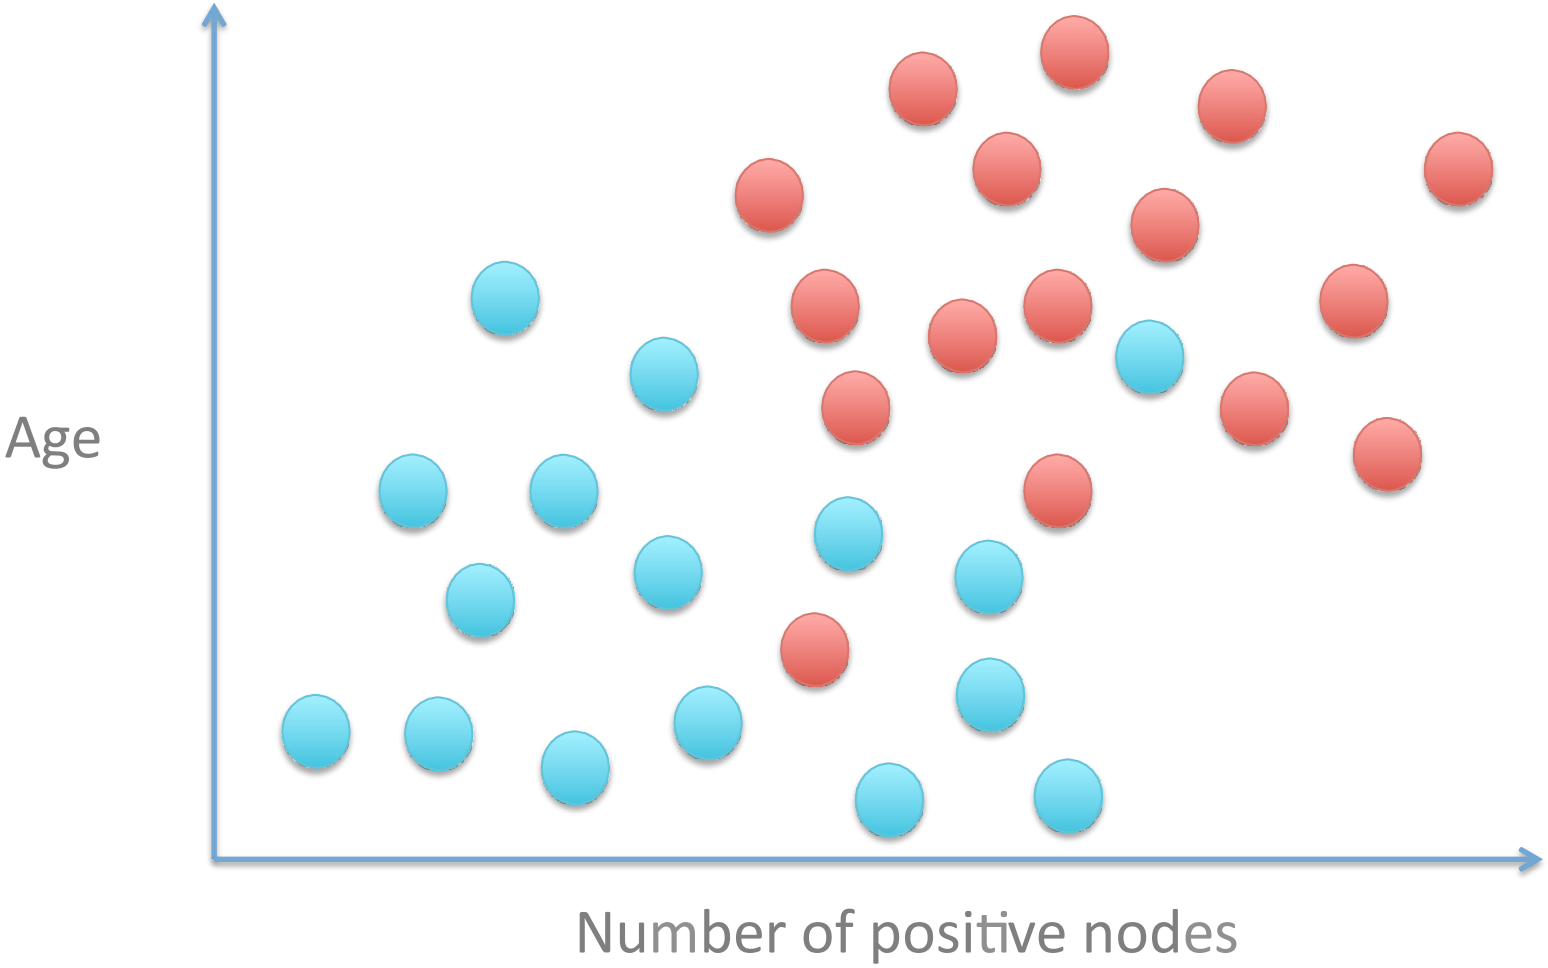

2 Features: Number of + nodes, Age

2 Labels: Survived / Lost

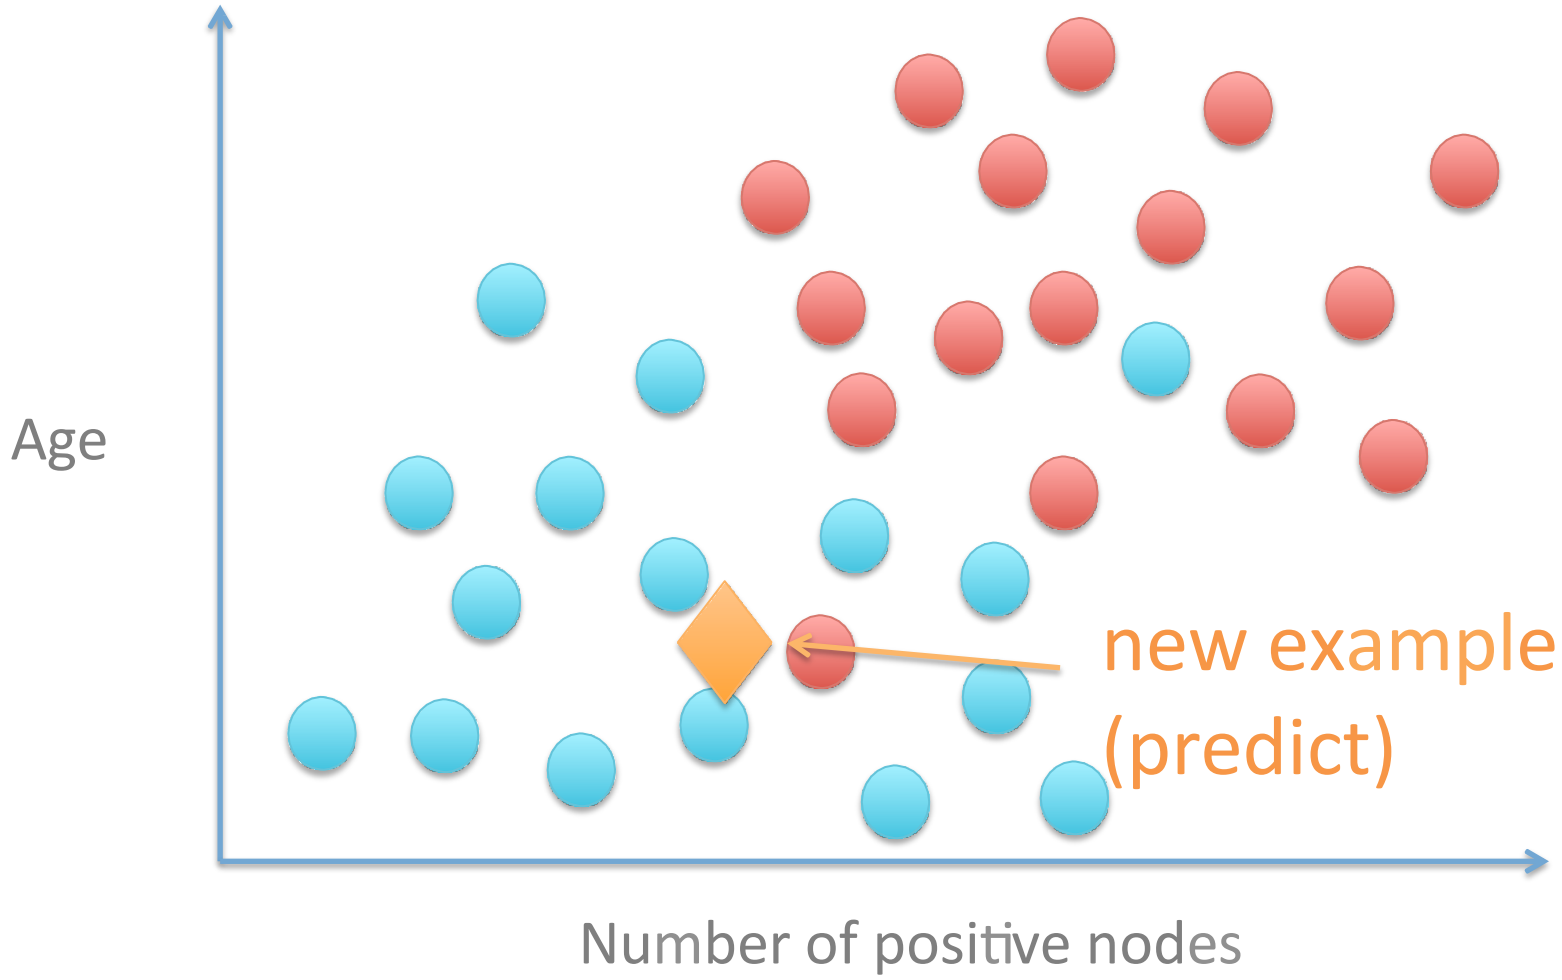

Binary classification:

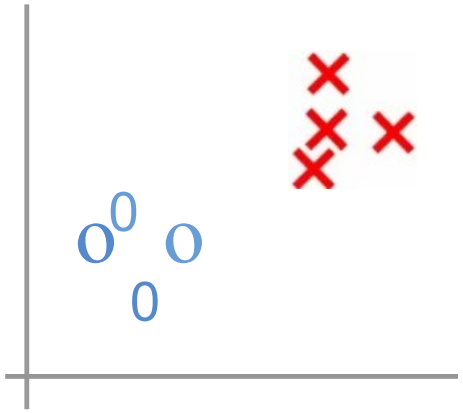

Multi-class classification:

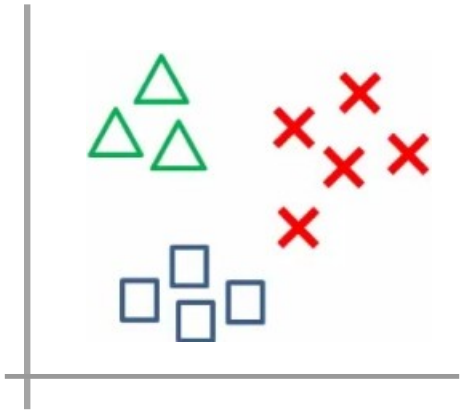

Linear regression for classification?

Linear regression for classification?

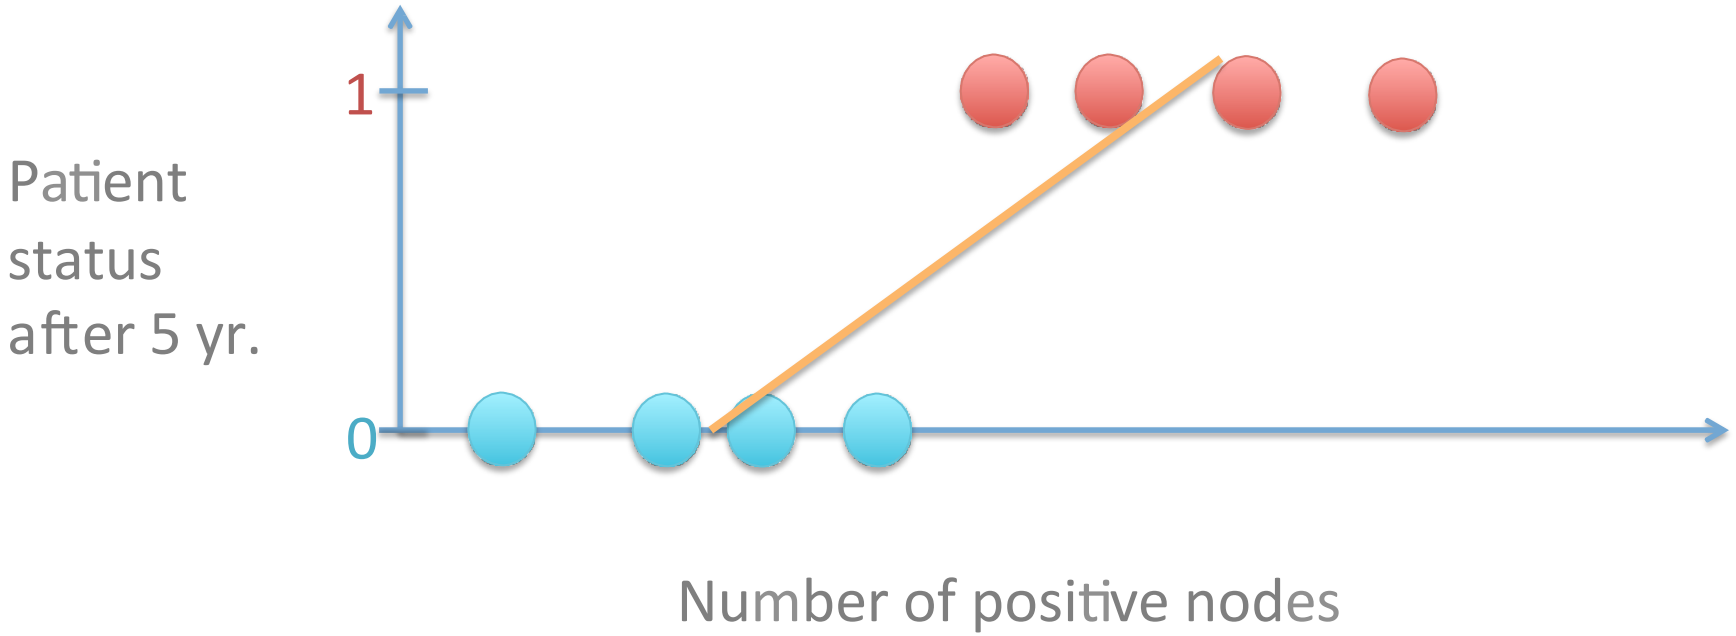

$$y_\beta(x) = \beta_0 + \beta_1 x + \varepsilon$$

Linear regression for classification?

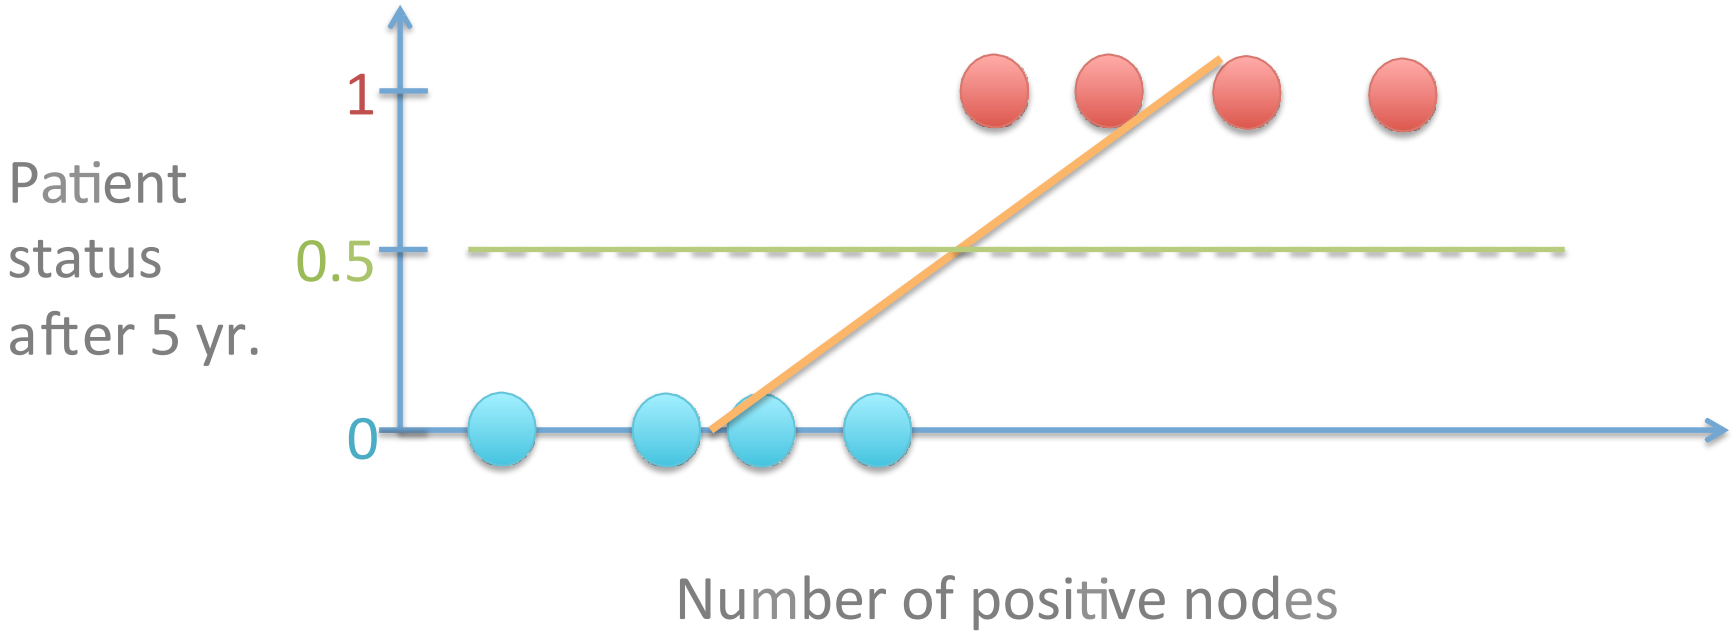

If y_pred > 0.5 predict label 1 (lost)
If y_pred < 0.5 predict label 0 (survived)

Linear regression for classification?

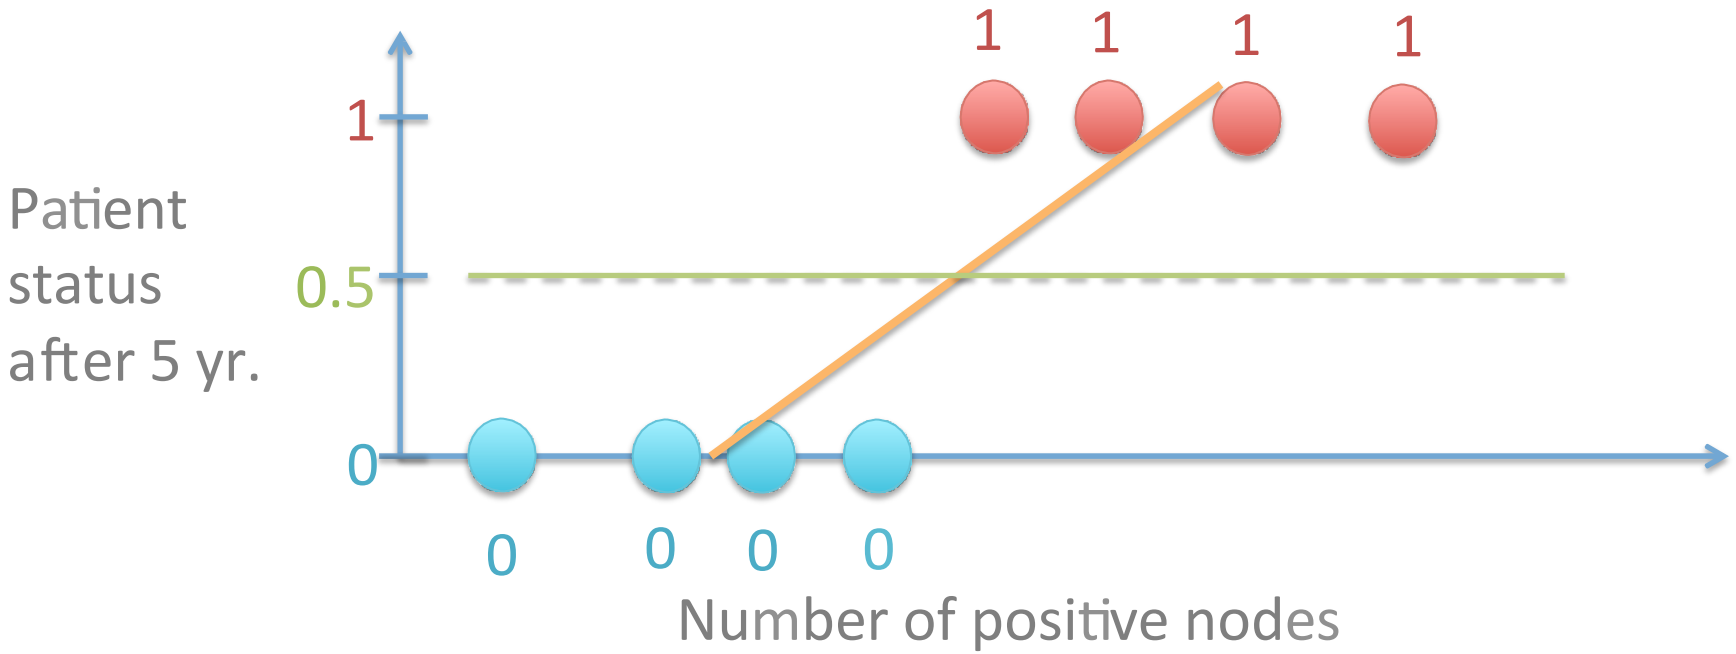

If y_pred > 0.5 predict label 1 (lost)
If y_pred < 0.5 predict label 0 (survived)

Linear regression for classification?

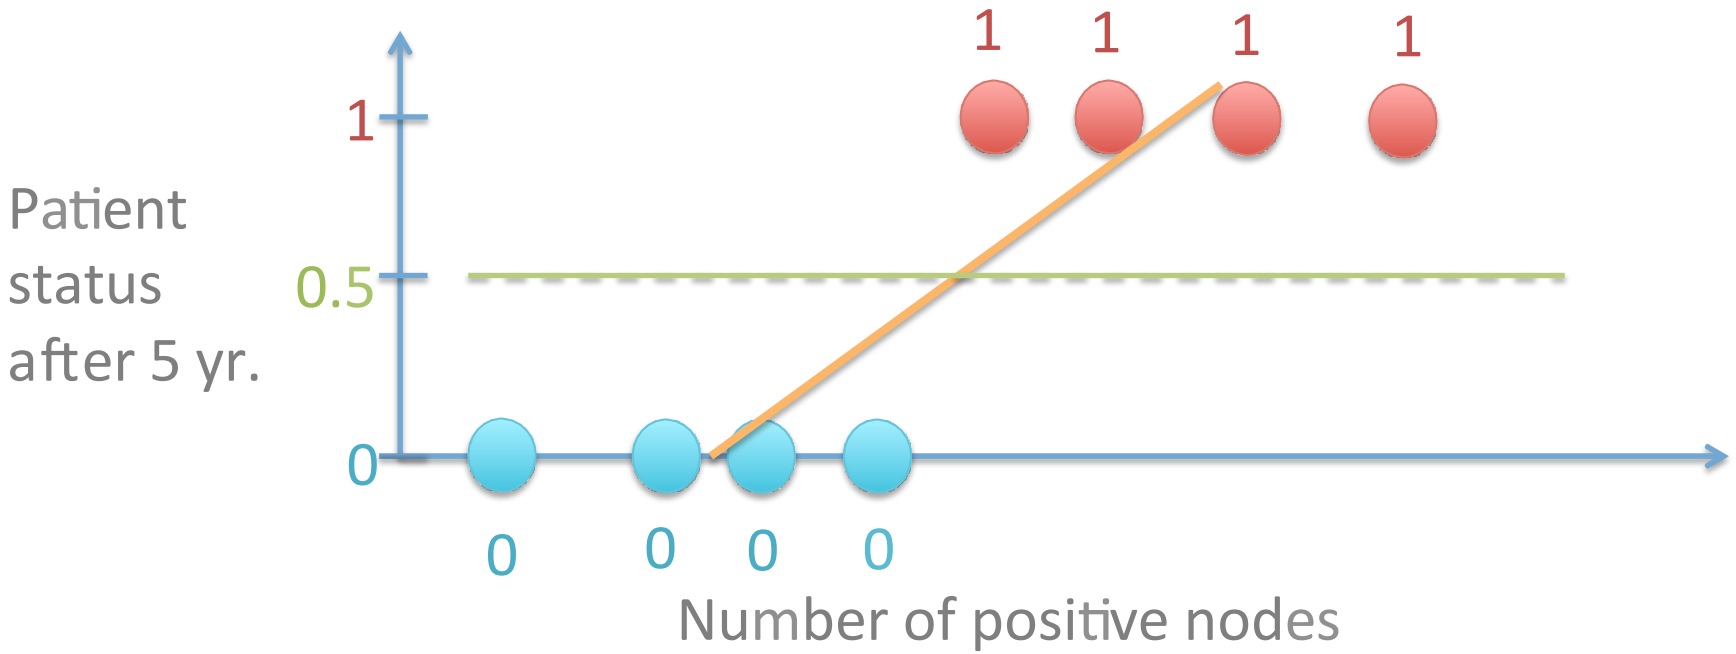

If y_pred > 0.5 predict label 1 (lost)
If y_pred < 0.5 predict label 0 (survived)

Linear regression for classification?

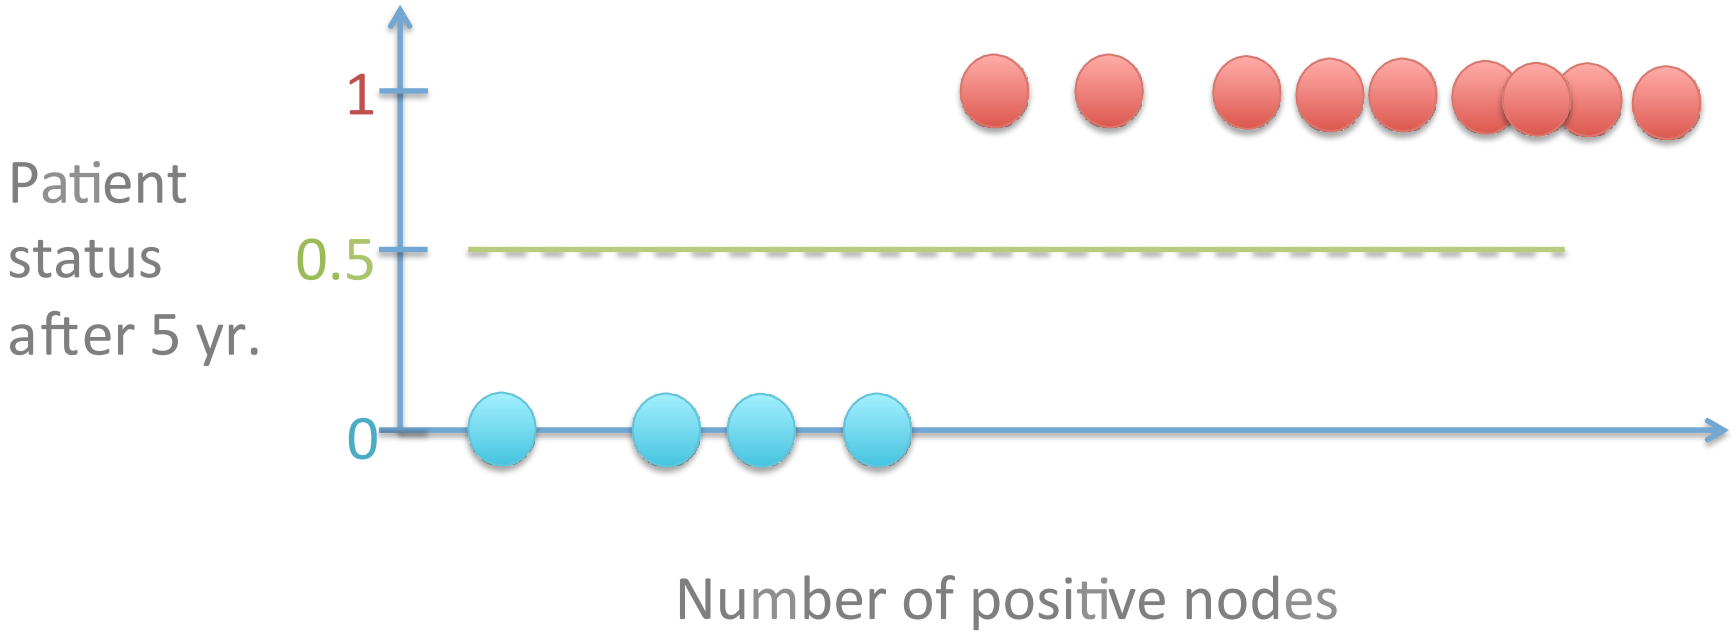

If y_pred > 0.5 predict label 1 (lost)
If y_pred < 0.5 predict label 0 (survived)

Linear regression for classification?

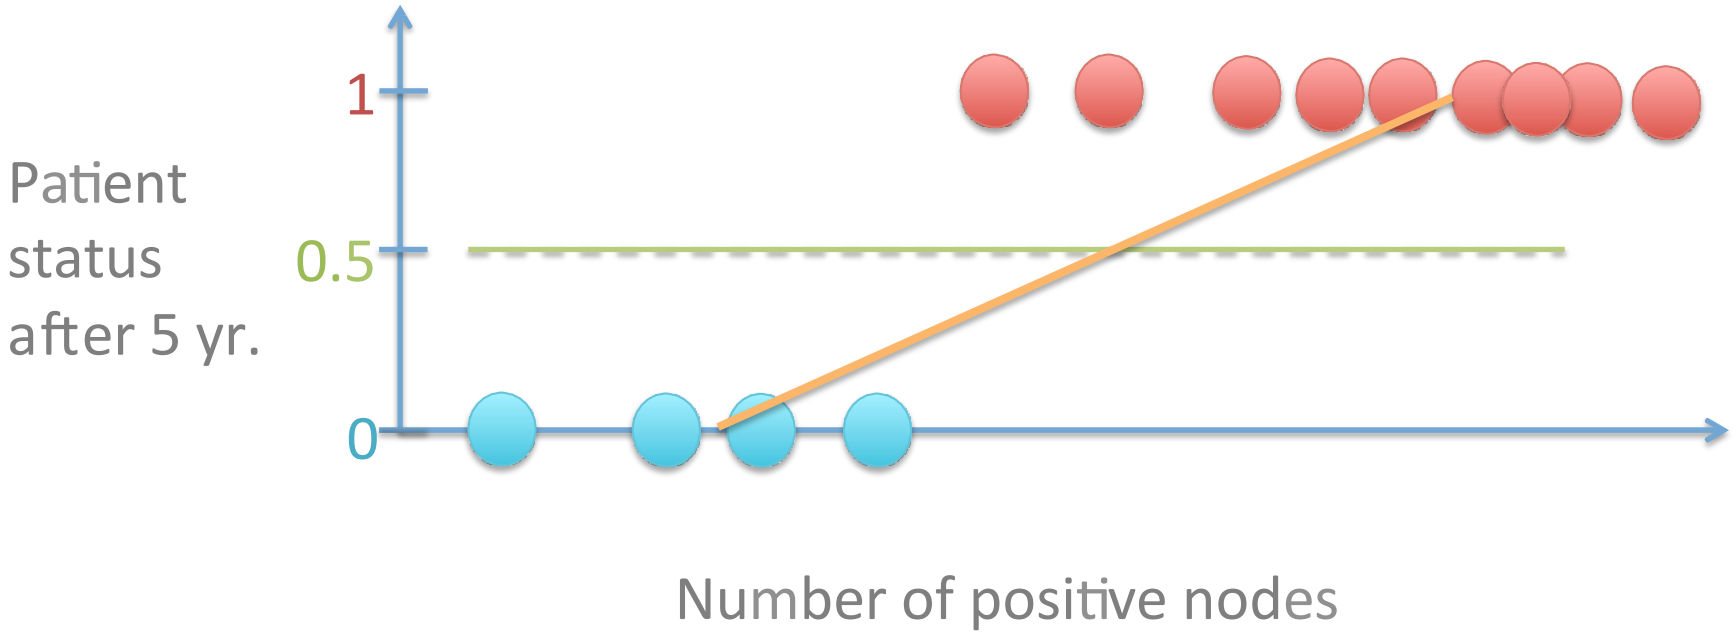

If y_pred > 0.5 predict label 1 (lost)
If y_pred < 0.5 predict label 0 (survived)

Linear regression for classification?

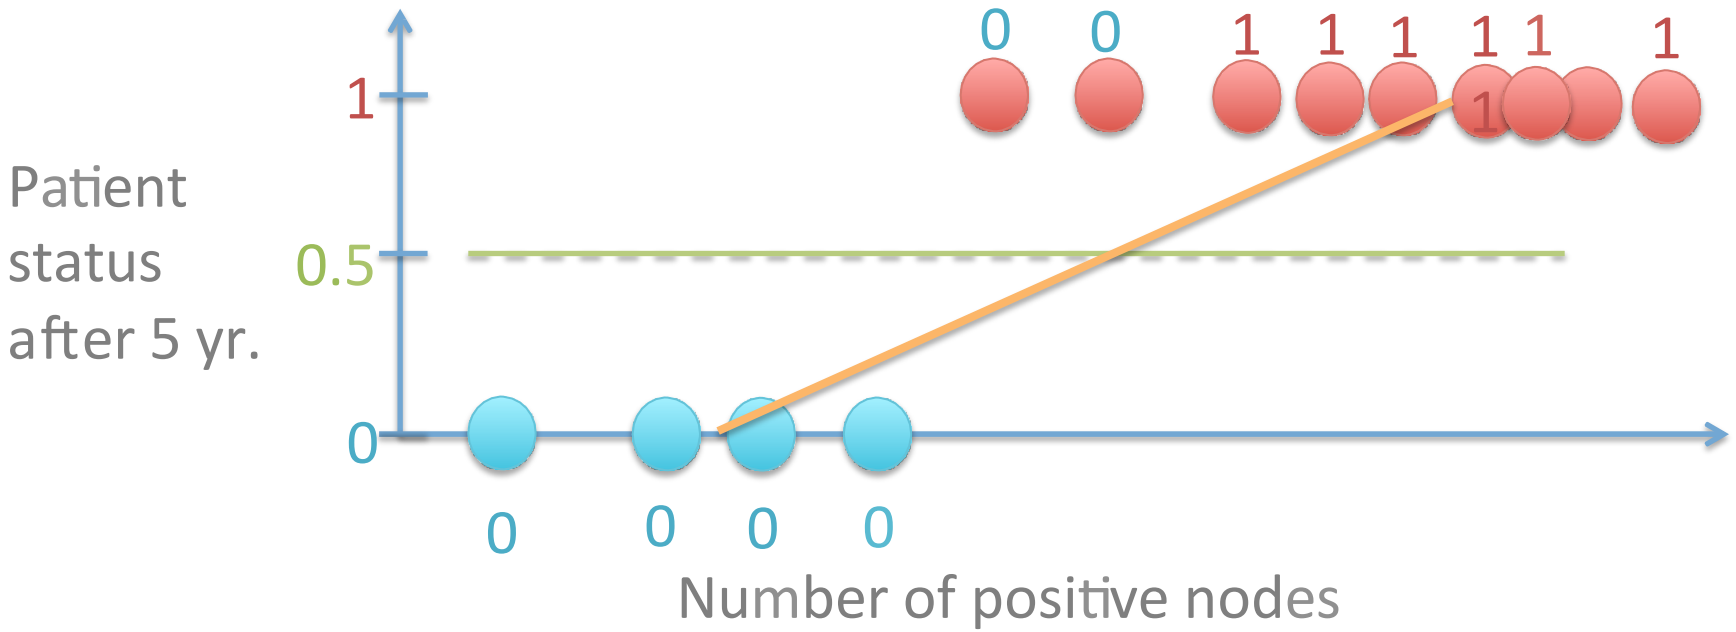

If y_pred > 0.5 predict label 1 (lost)
If y_pred < 0.5 predict label 0 (survived)

Linear regression for classification?

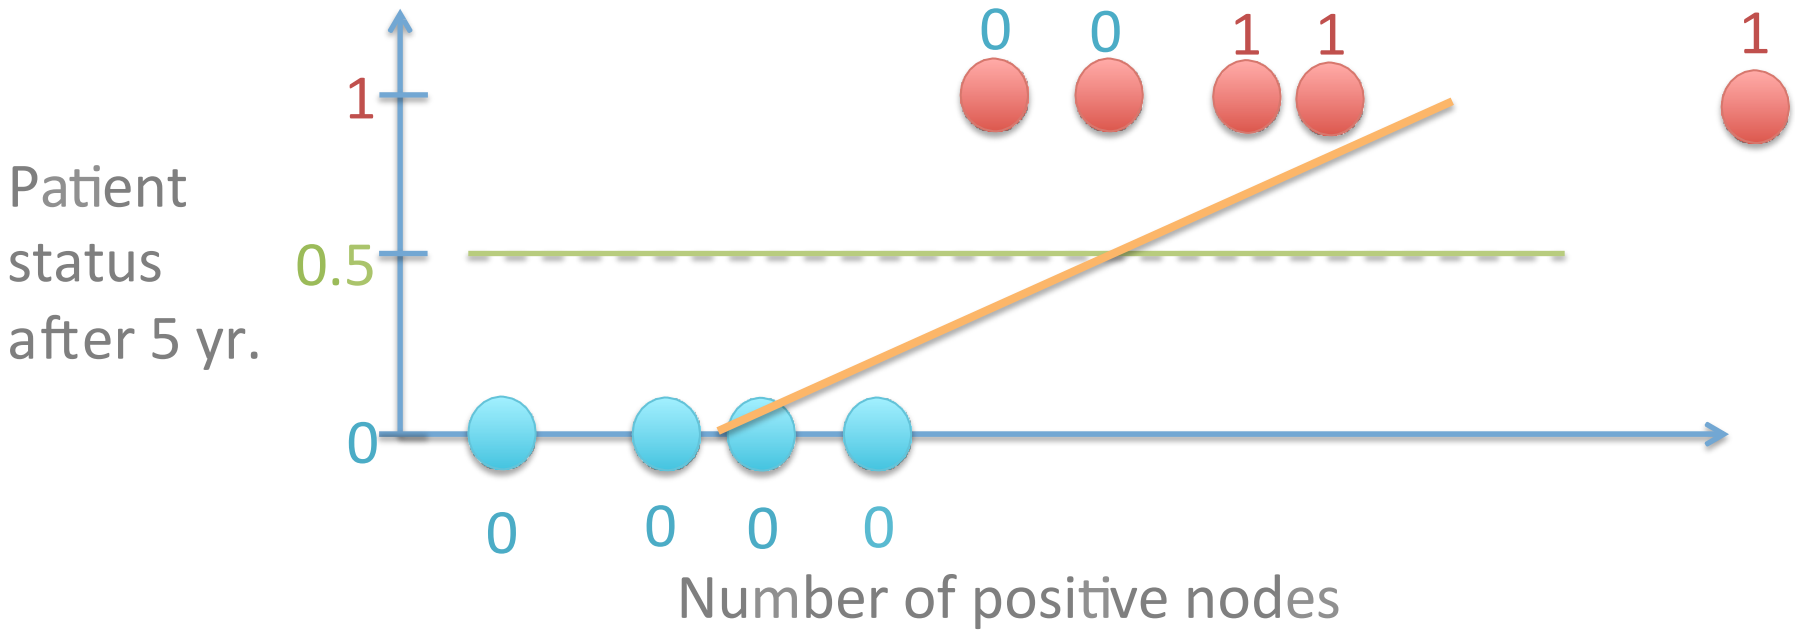

If y_pred > 0.5 predict label 1 (lost)
If y_pred < 0.5 predict label 0 (survived)

Logistic regression to the rescue

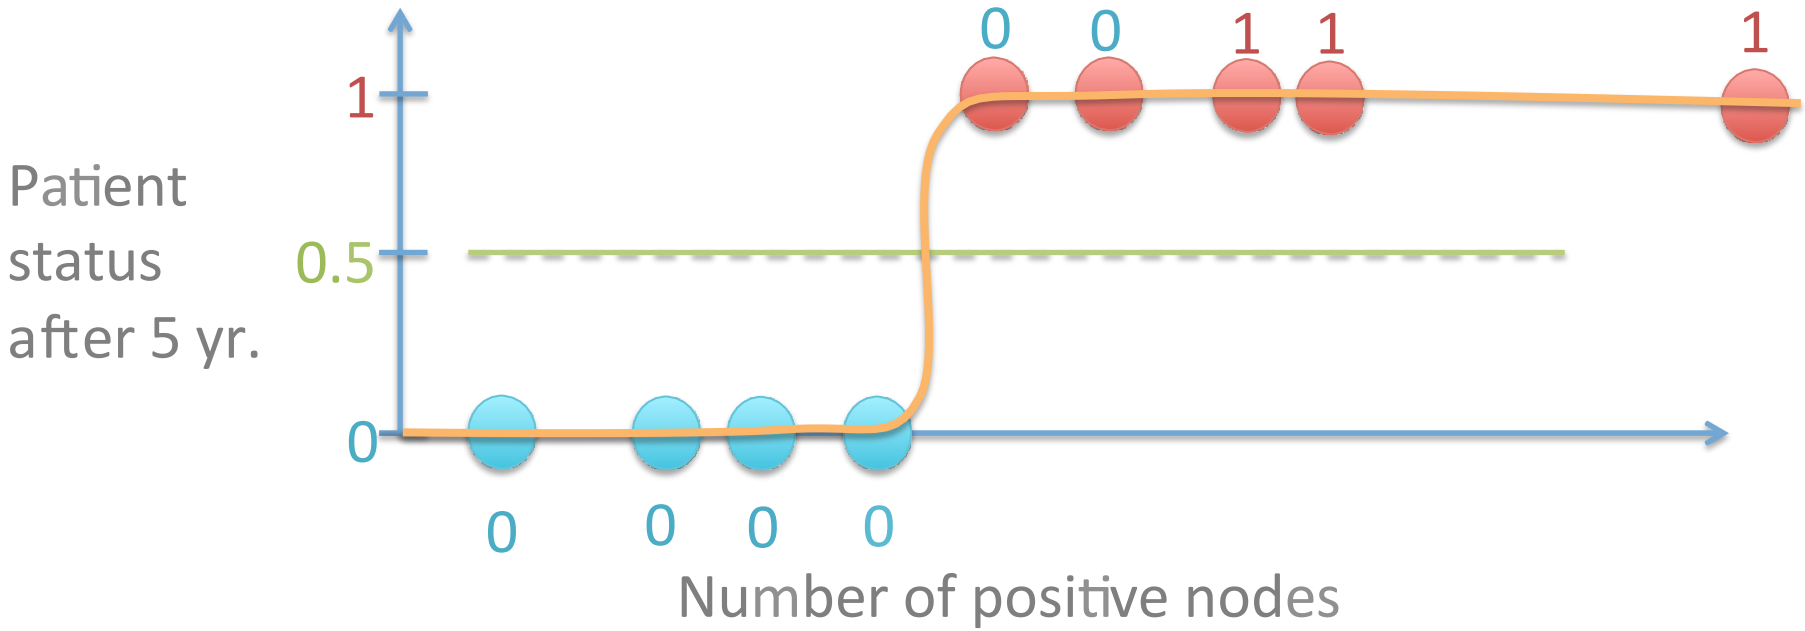

$$y_\beta(x) = \frac{1}{1+e^{-(\beta_0+\beta_1 x+\varepsilon)}}$$

What is this function?

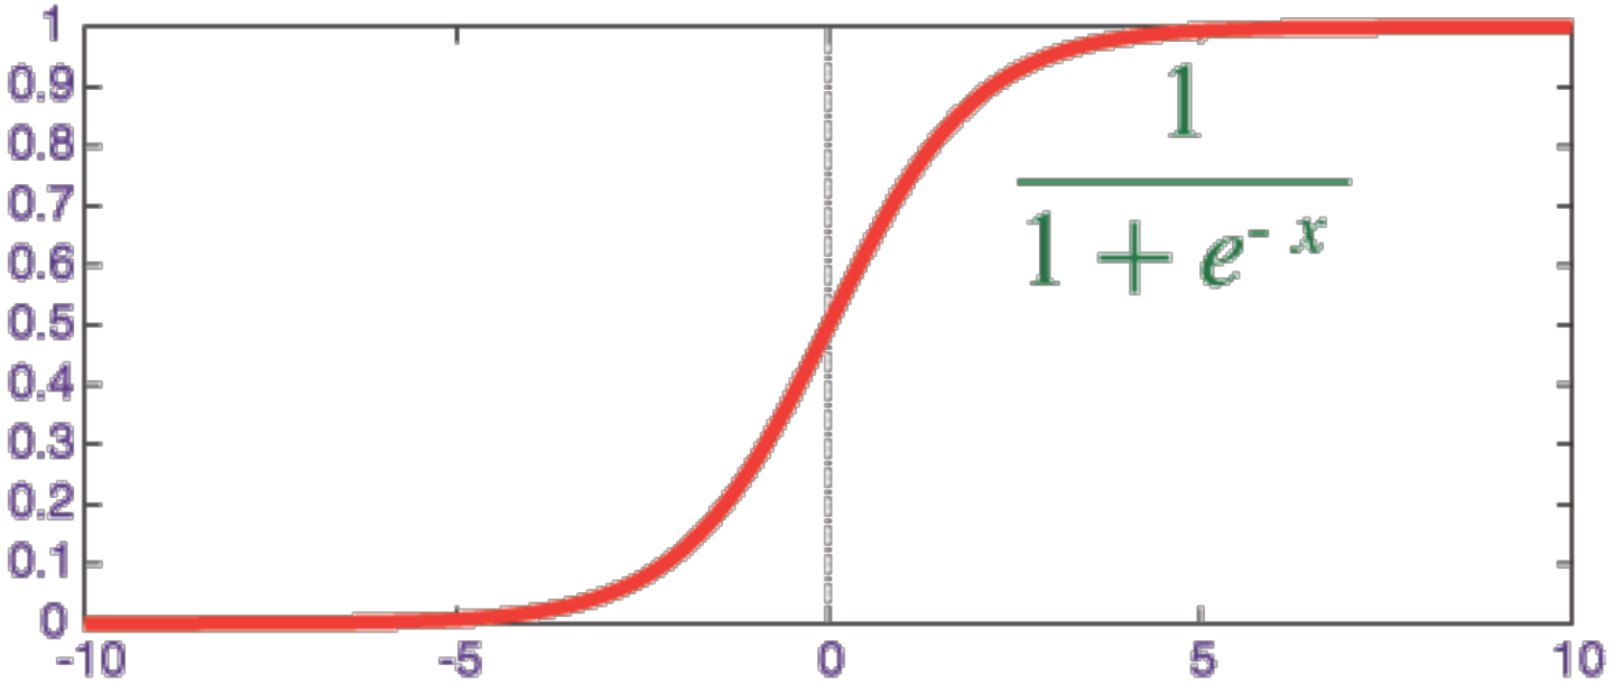

What is this function?

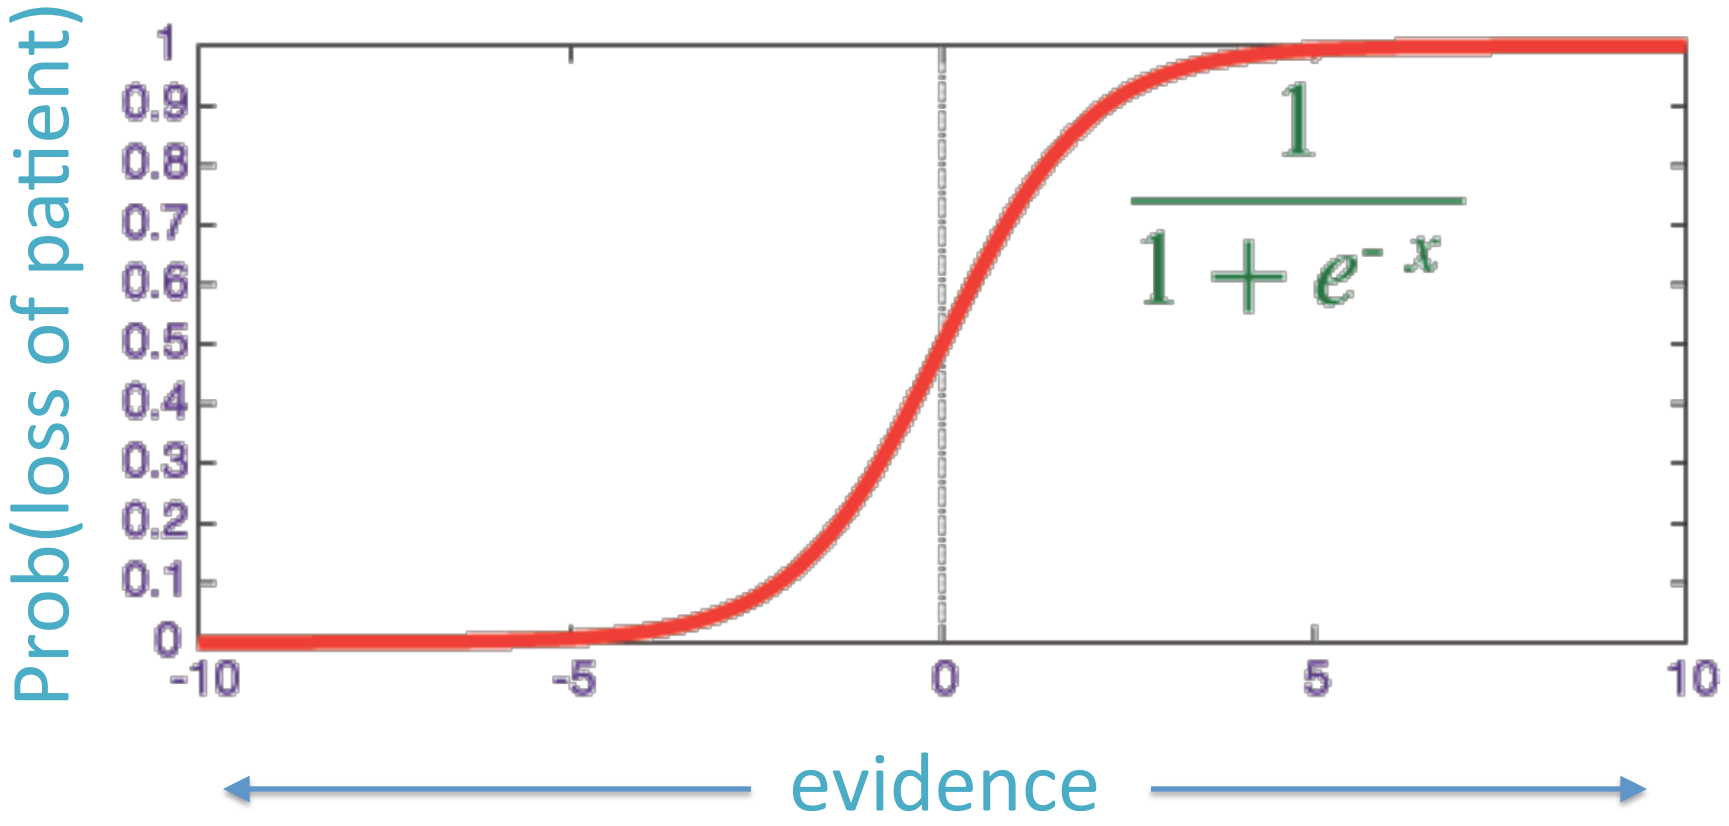

What is this function?

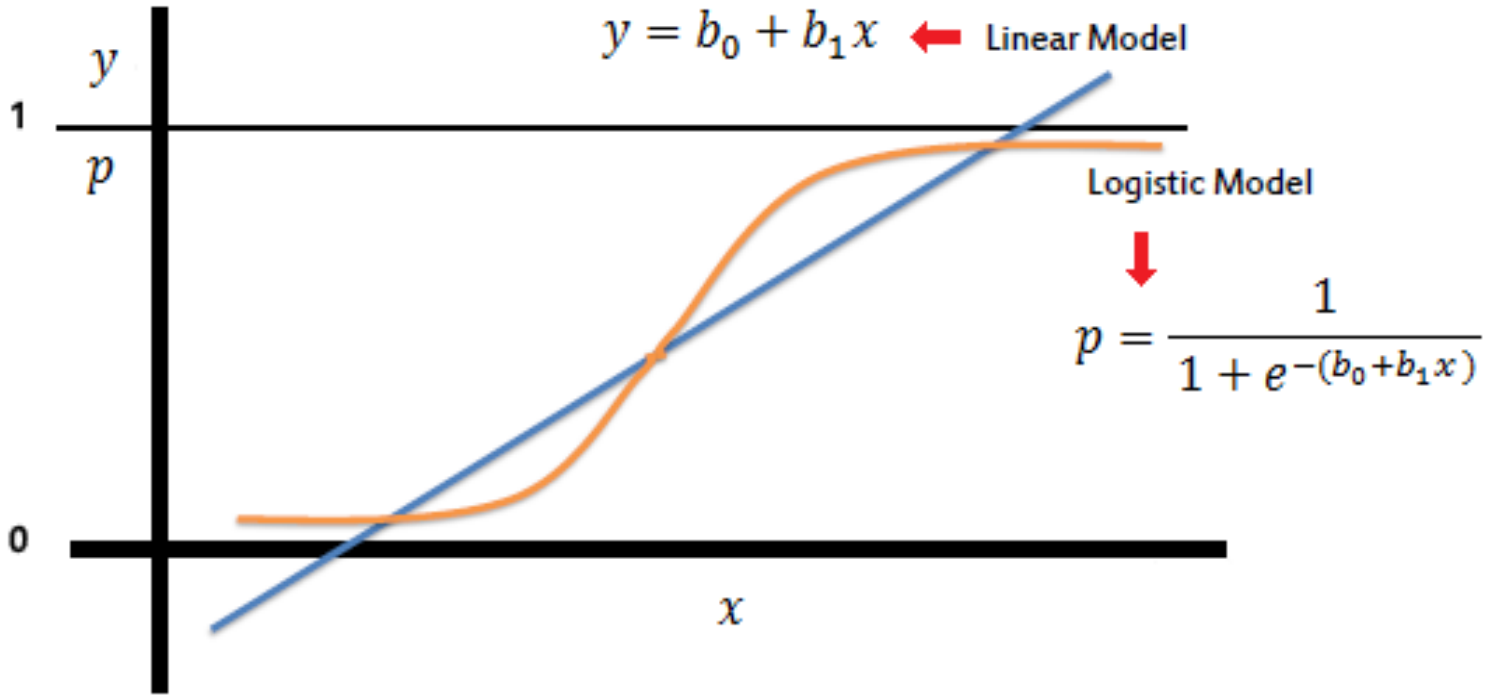

http://www.saedsayad.com/logistic_regression.htm

Linear regression for classification?

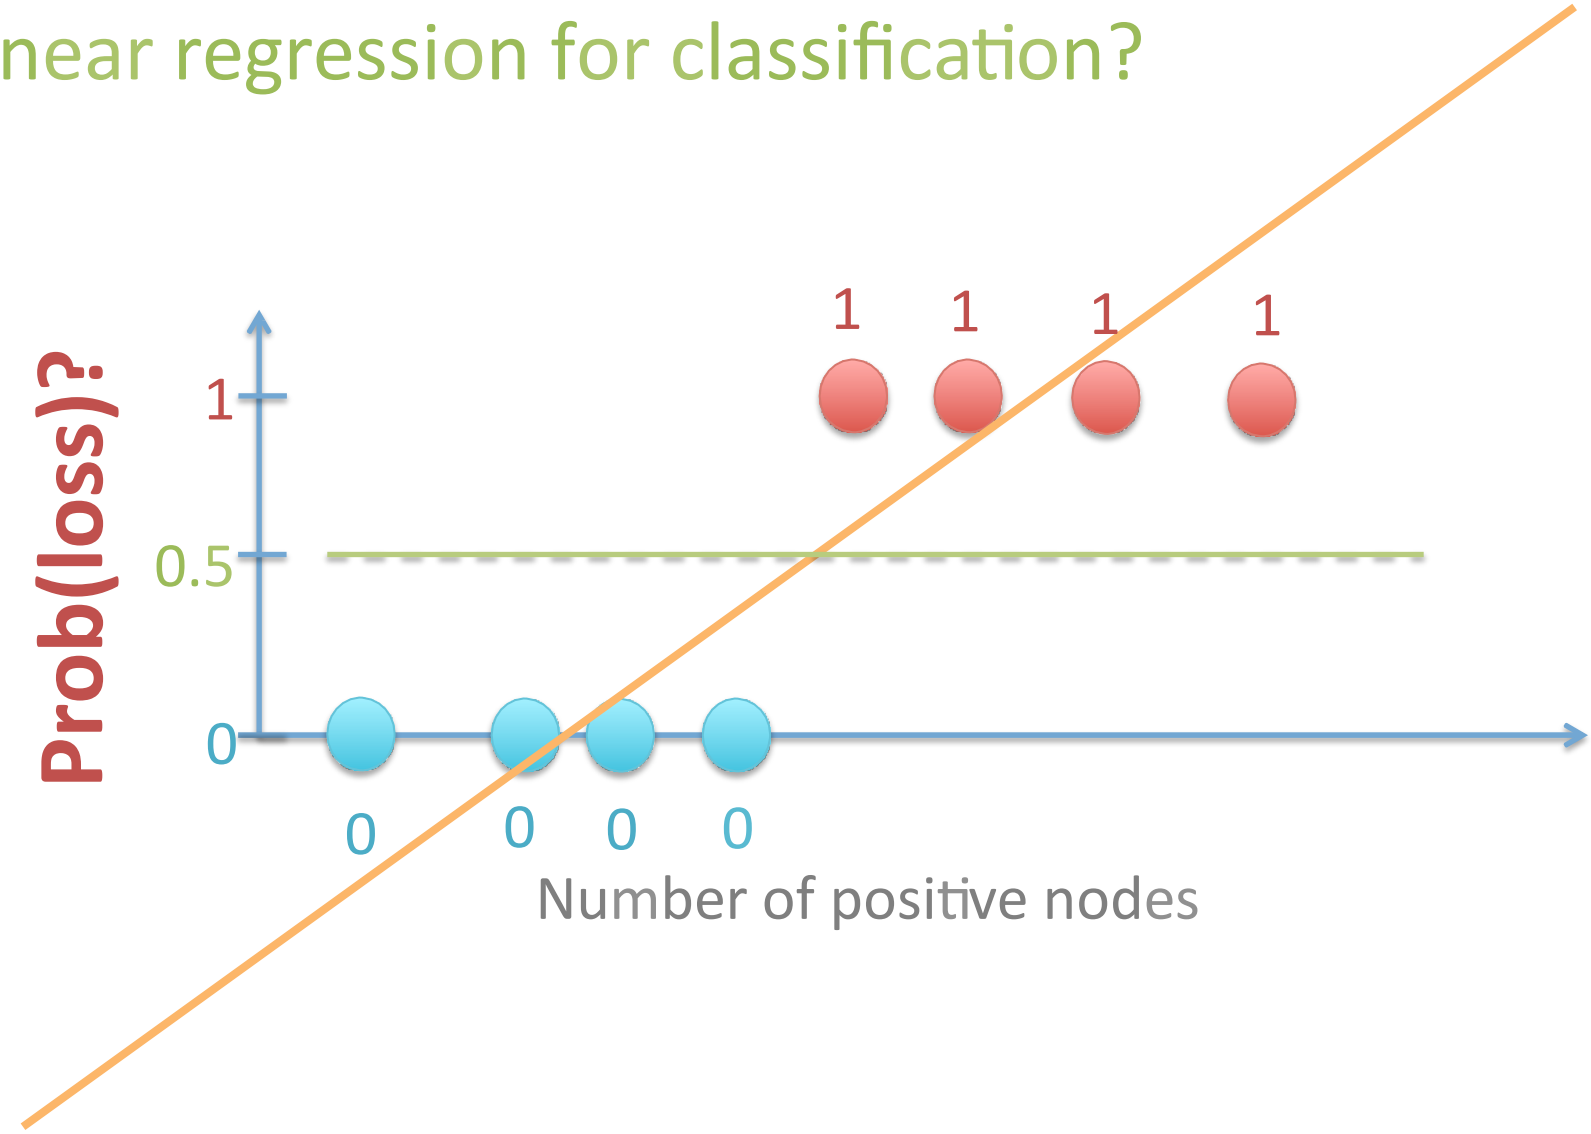

Linear regression for classification?

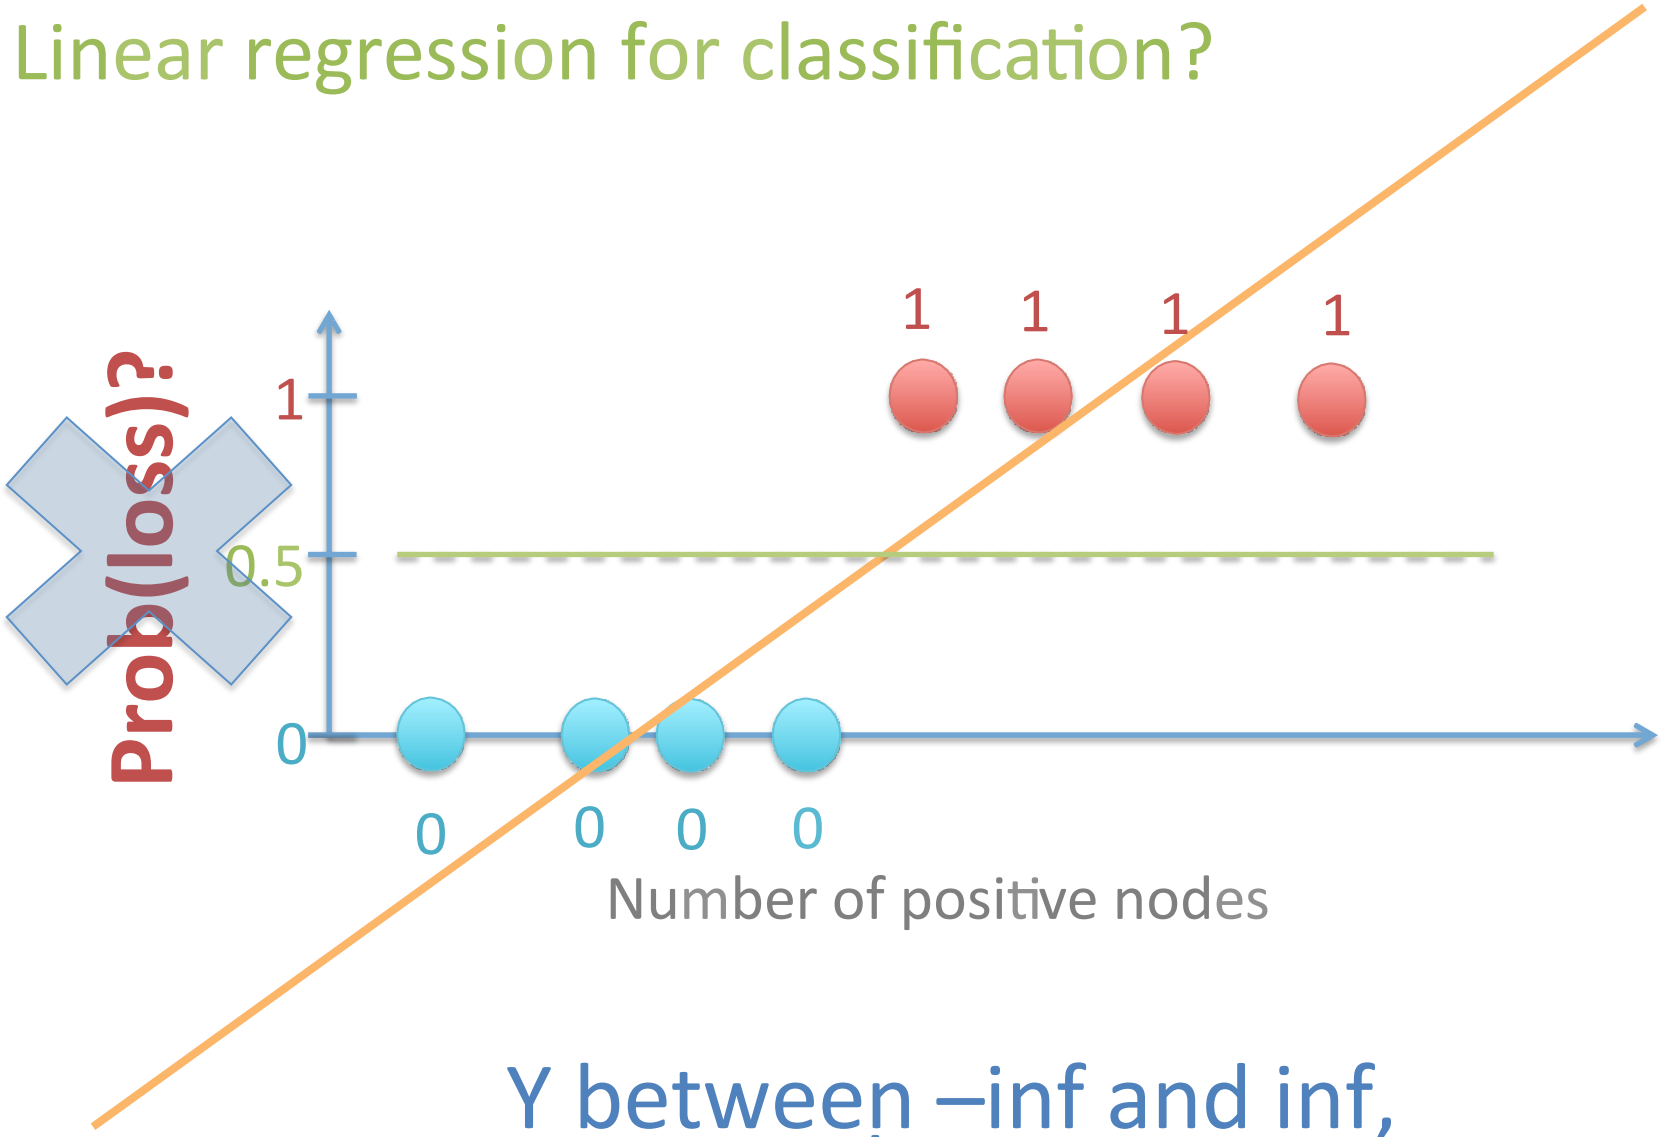

Y between –inf and inf,
Not 0 and 1

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.8$$

$$P(survival) = 0.2$$

Probability

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.8$$

$$P(survival) = 0.2$$

Probability

$$\frac{P(loss)}{P(survival)} = 4$$

Odds

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.05$$
$$P(survival) = 0.95$$

Probability

$$\frac{P(loss)}{P(survival)} = 0.053$$

Odds

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.5$$

$$P(survival) = 0.5$$

Probability

$$\frac{P(loss)}{P(survival)} = 1$$

Odds

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.5$$

$$P(survival) = 0.5$$

Probability

$$\frac{P(loss)}{P(survival)} = 1$$

Odds

between 0 and inf

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.5$$

$$P(survival) = 0.5$$

Probability

$$\log\left(\frac{P(loss)}{P(survival)}\right) = 0$$

Log Odds

between -inf and inf

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.05$$
$$P(survival) = 0.95$$

Probability

$$\log\left(\frac{P(loss)}{P(survival)}\right) = -2.94$$

Log Odds

between -inf and inf

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.8$$
$$P(survival) = 0.2$$

Probability

$$\log\left(\frac{P(loss)}{P(survival)}\right) = 1.39$$

Log Odds

between -inf and inf

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.999$$
$$P(survival) = 0.001$$

Probability

$$\log\left(\frac{P(loss)}{P(survival)}\right) = 6.9$$

Log Odds

between -inf and inf

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.999$$

$$1 - P(loss) = 0.001$$

Probability

$$\log\left(\frac{P(loss)}{1 - P(loss)}\right) = 6.9$$

Log Odds

logit function

What metric would express the chances of loss/survival, but not constrained to [0, 1] ?

$$P(loss) = 0.999$$
$$1 - P(loss) = 0.001$$

Probability

$$\log\left(\frac{P(loss)}{1 - P(loss)}\right) = 6.9$$

Log Odds

logit function

$$\frac{1}{1 - e^{\log\left(\frac{P(loss)}{1-P(loss)}\right)}} = P(loss)$$

Logistic Function

Log Odds → Prob

What is this function?

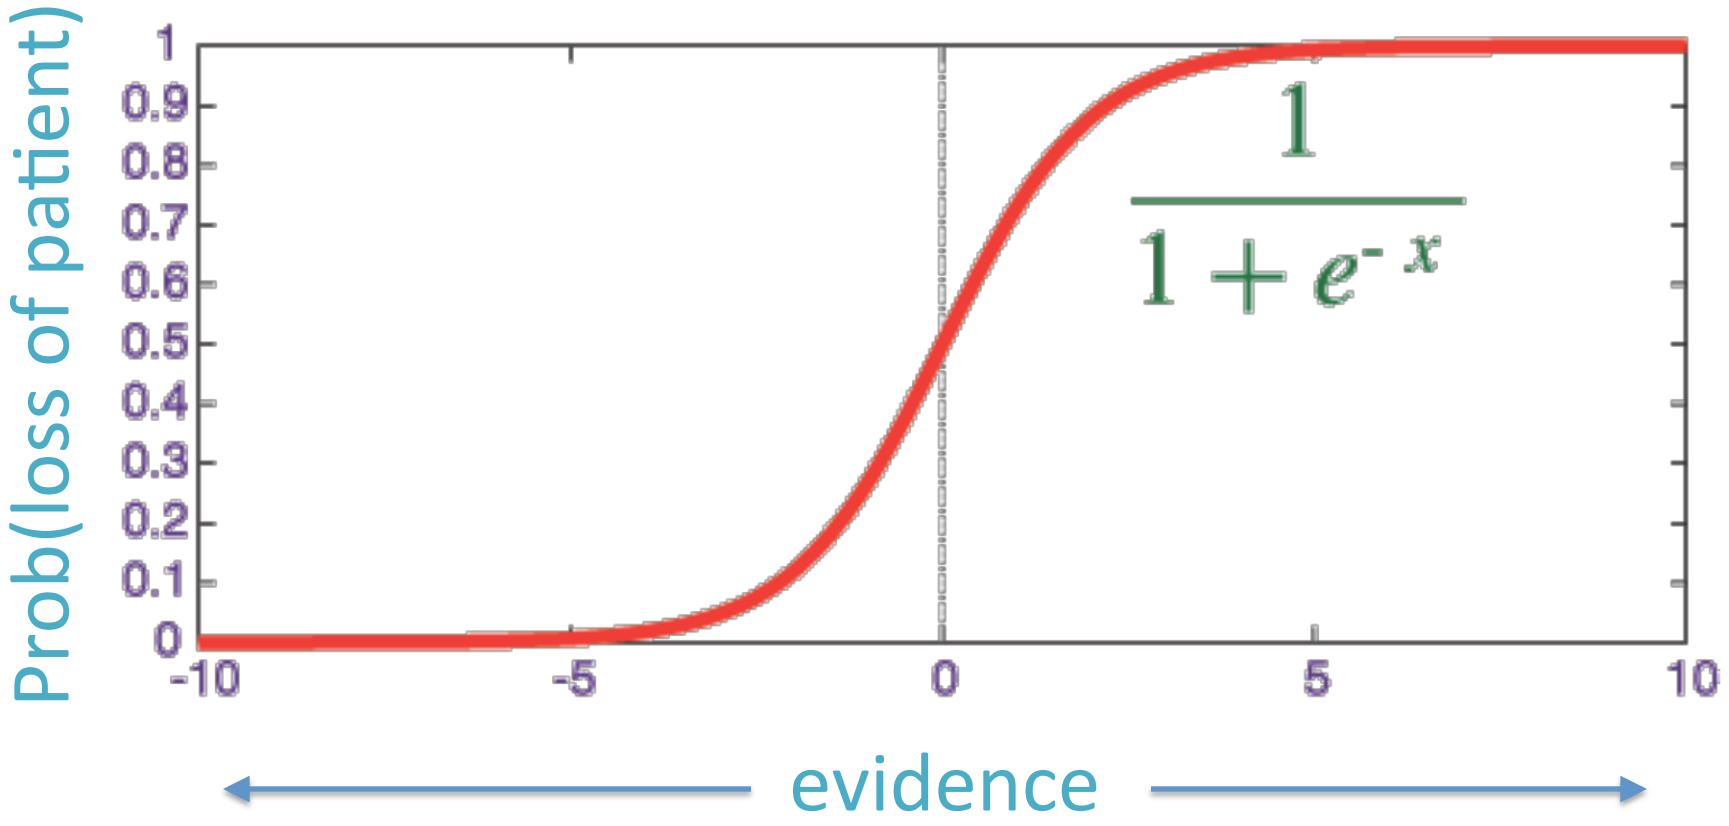

"Logistic regression" is a classification algorithm

$$y_\beta(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x + \varepsilon)}}$$

```
from sklearn.linear_model import LogisticRegression
#(just like LinearRegression)
```

```
from statsmodels.formula.api import Logit
#(just like OLS)
```

K Nearest Neighbors

K Nearest Neighbors

Number of positive nodes

K Nearest Neighbors

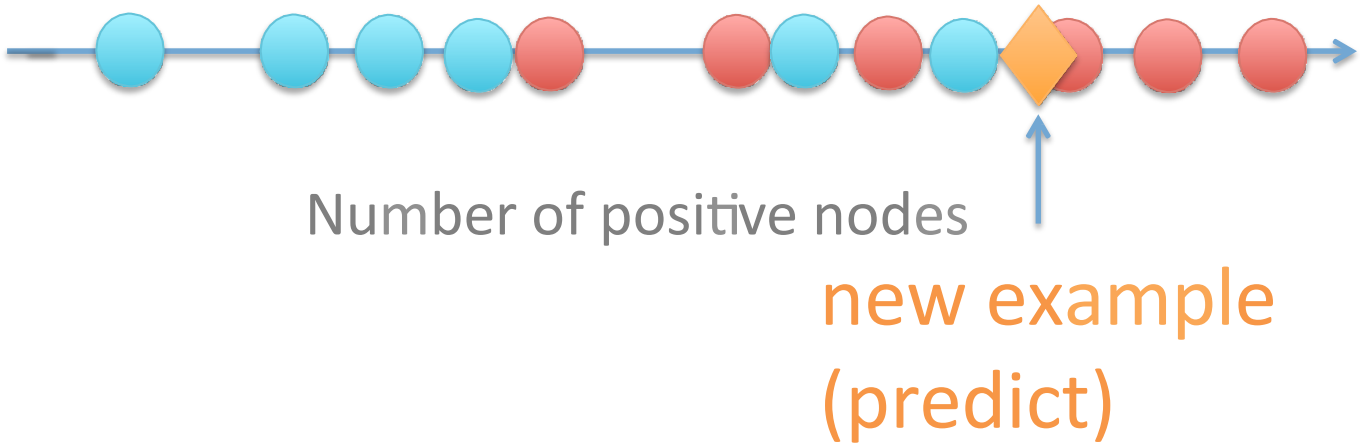

K Nearest Neighbors

K=1:

Look at the nearest neighbor, predict their label

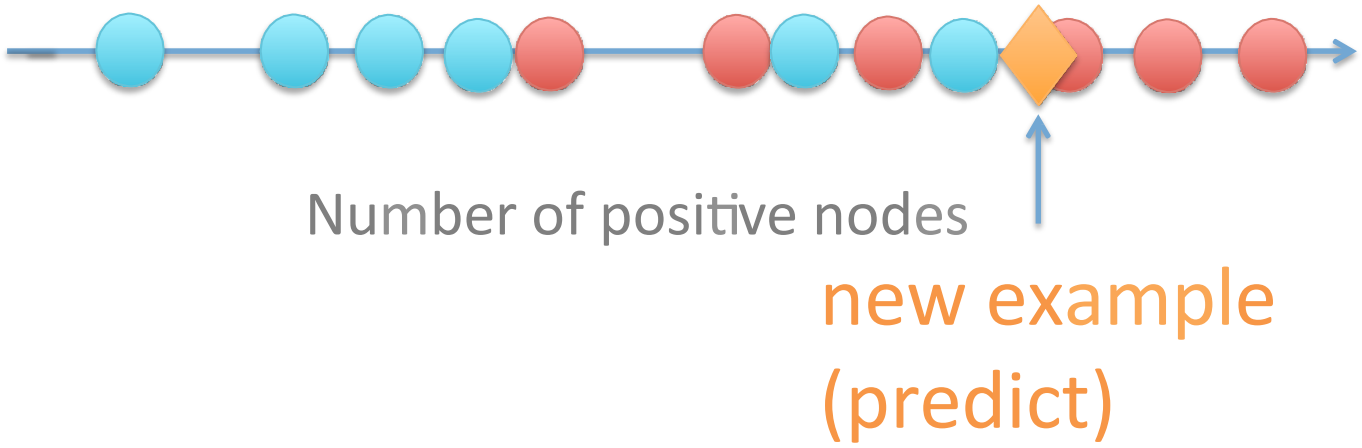

K Nearest Neighbors

K=1:

Look at the nearest neighbor, predict their label

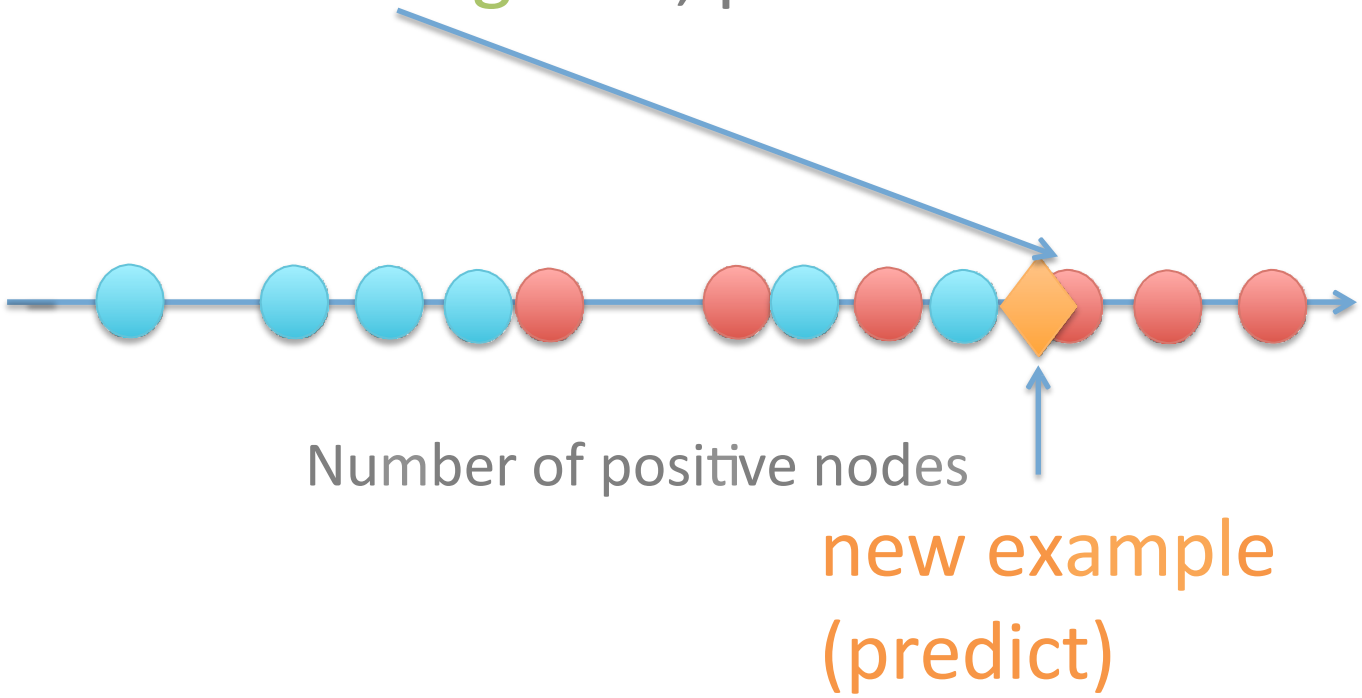

K Nearest Neighbors

K=1:

Look at the nearest neighbor, predict their label

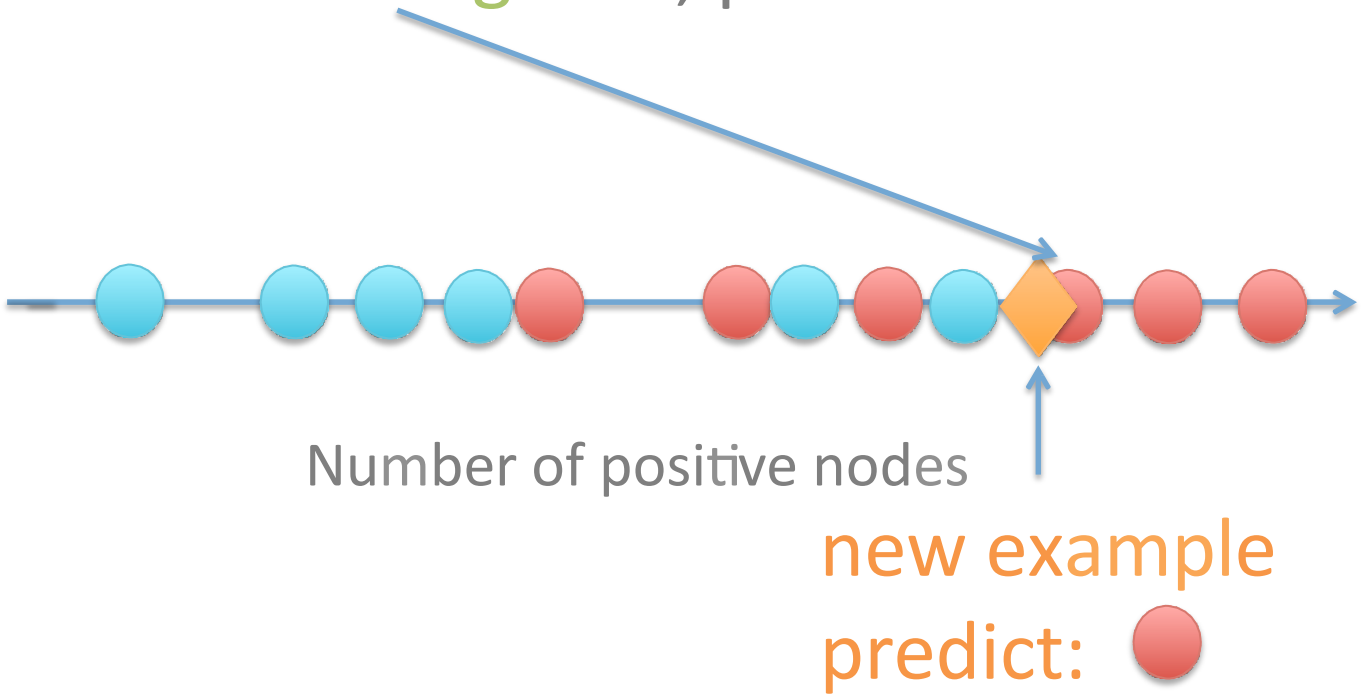

K Nearest Neighbors

K=2:

Look at the 2 nearest neighbors, predict the label you see the most

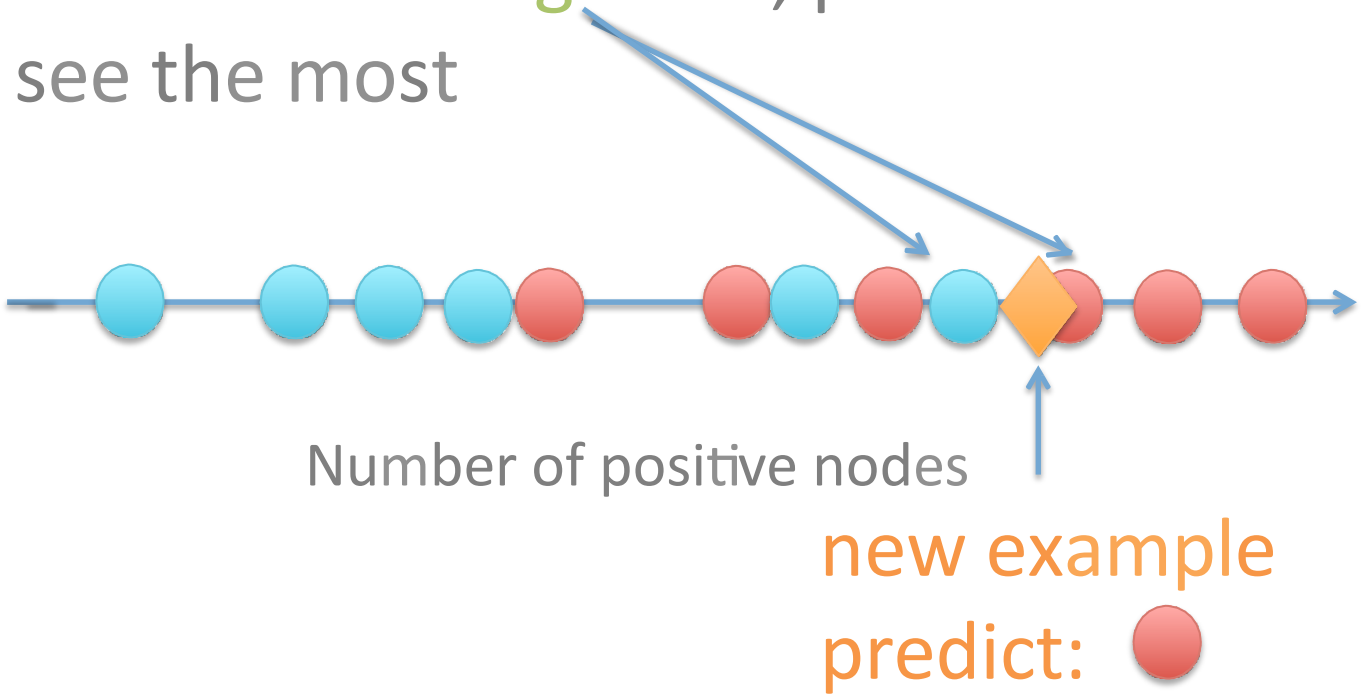

K Nearest Neighbors

K=3:

Look at the 3 nearest neighbors, predict the label you see the most

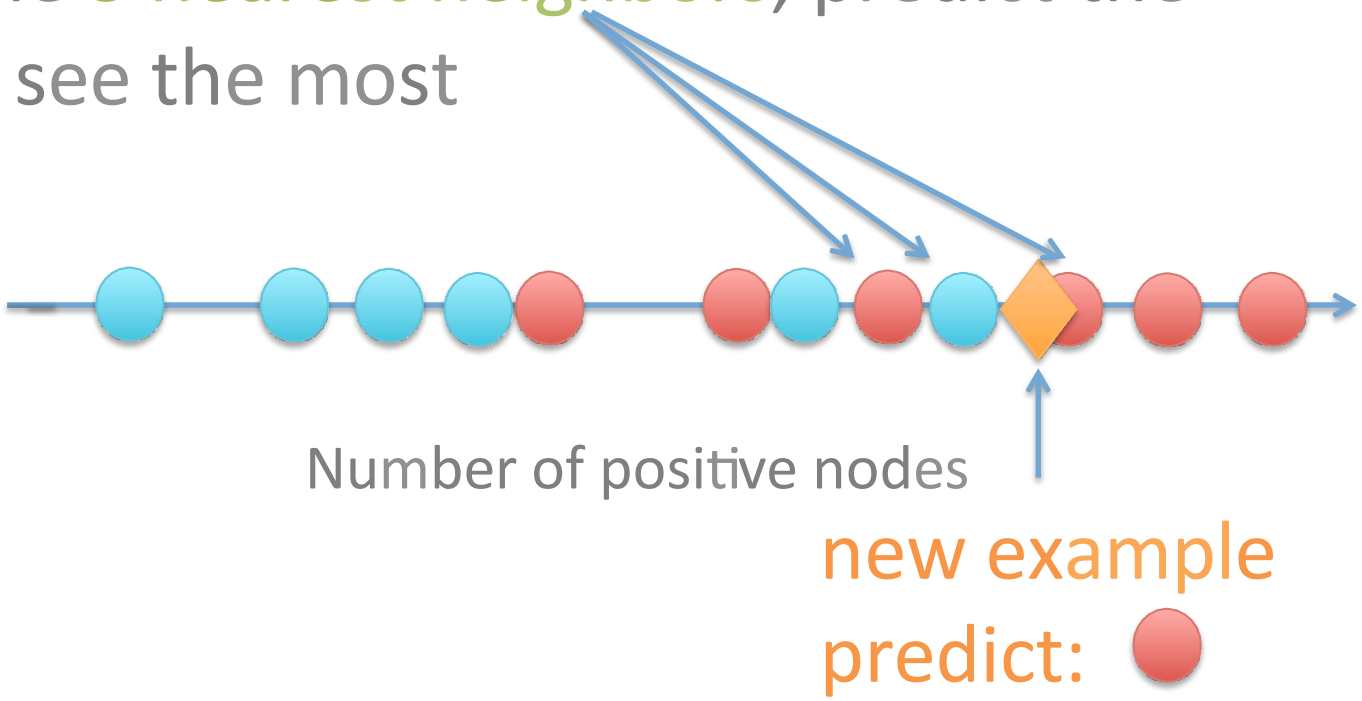

K Nearest Neighbors

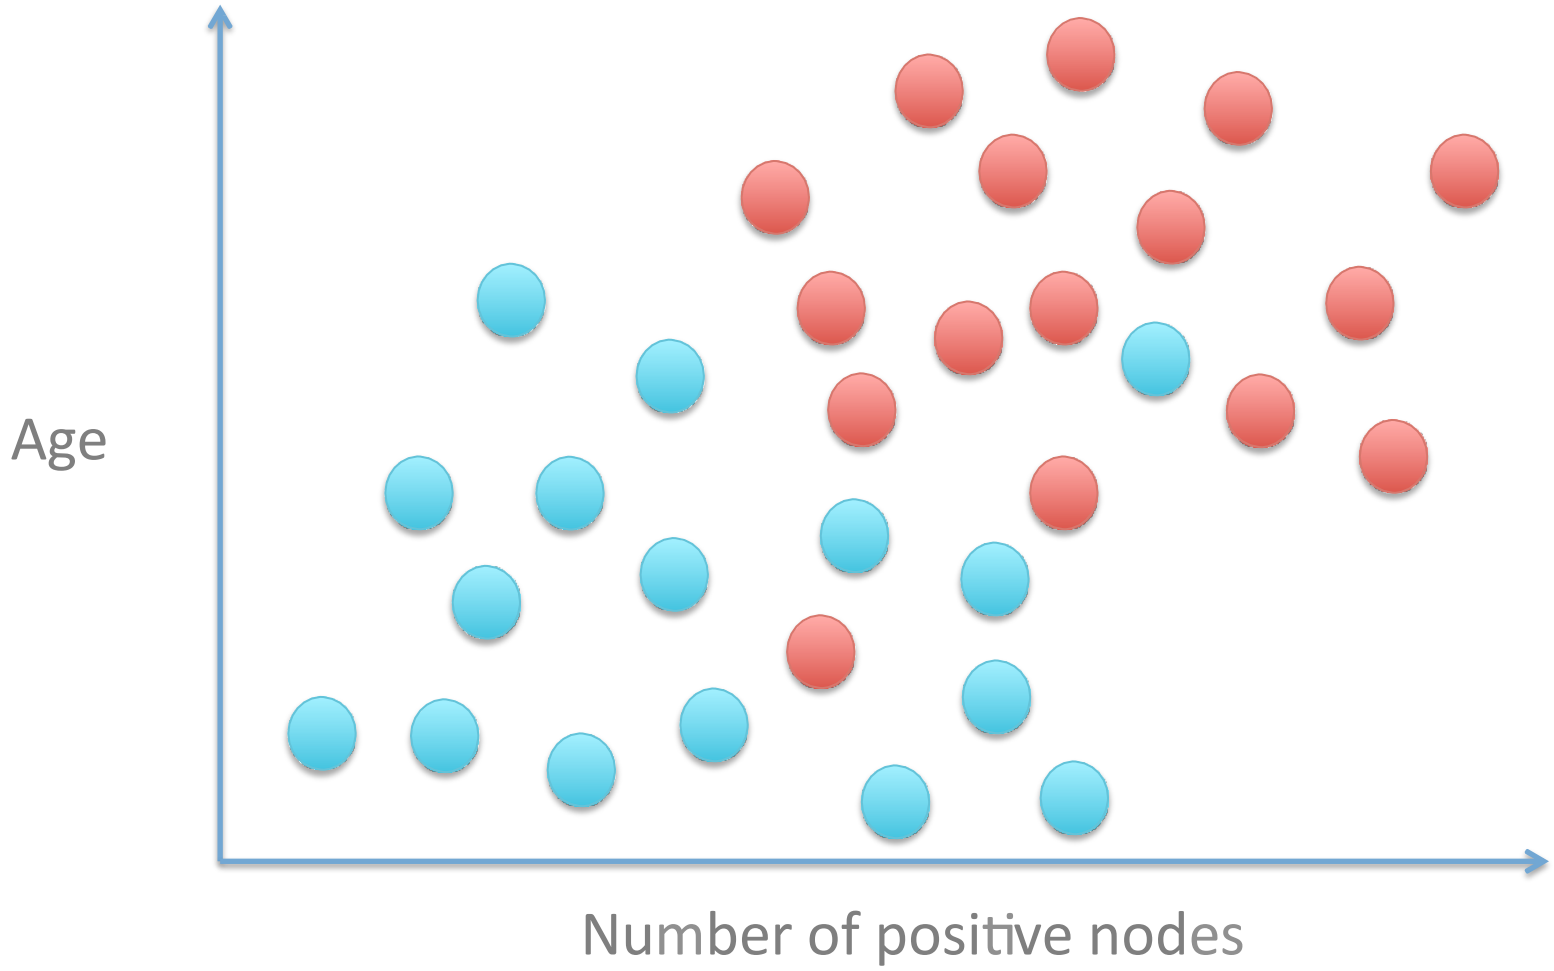

K Nearest Neighbors

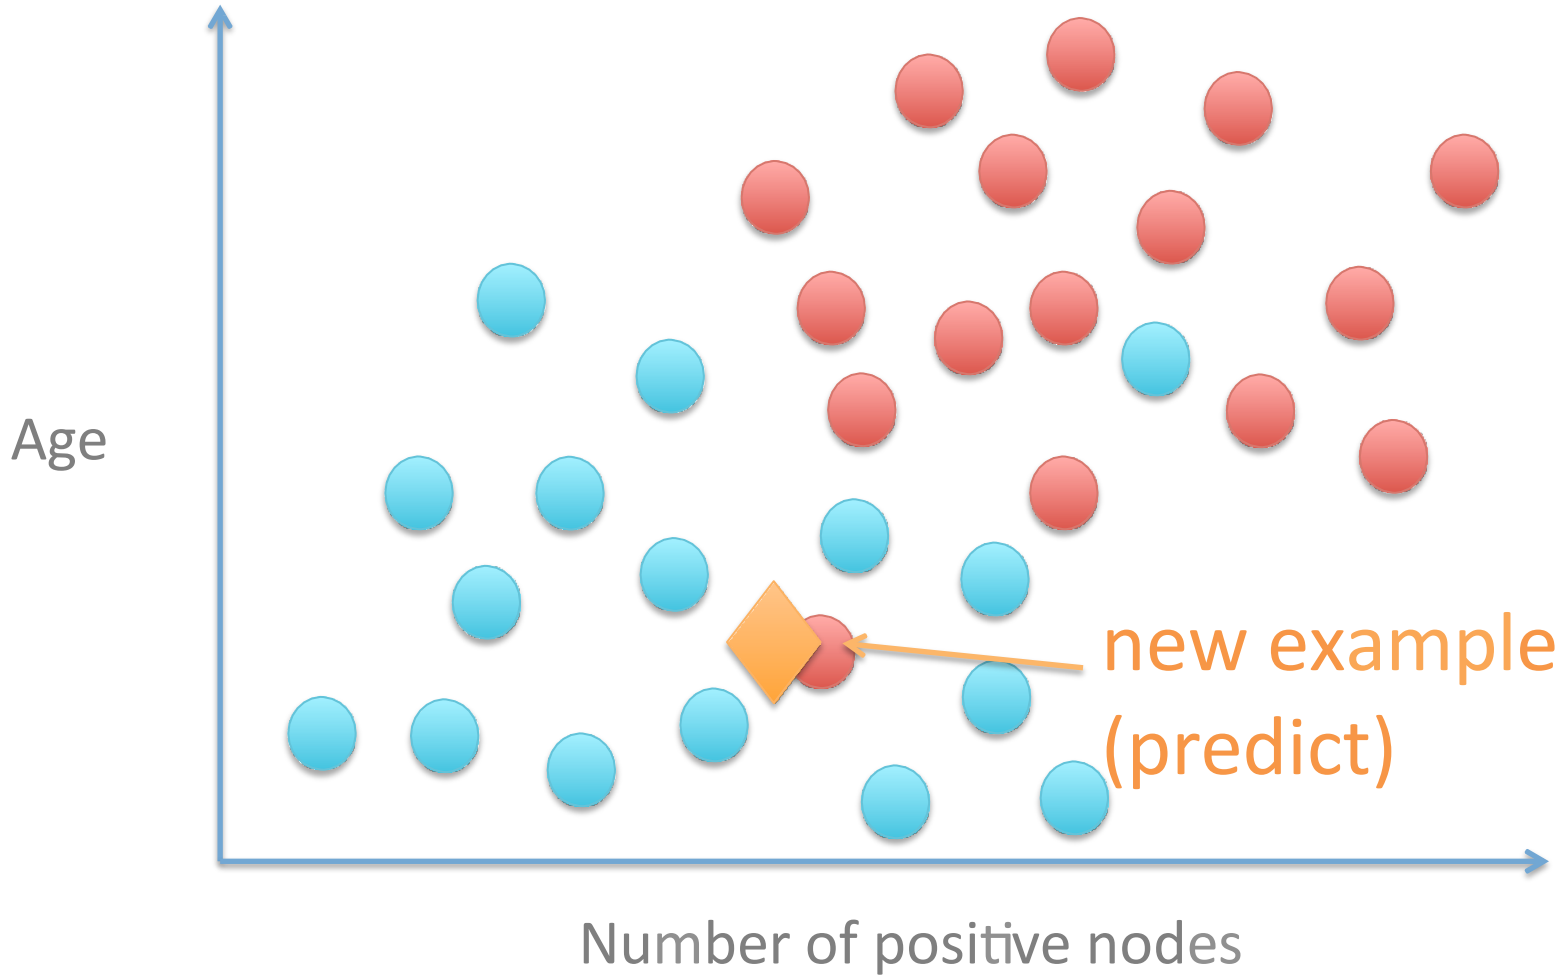

K Nearest Neighbors

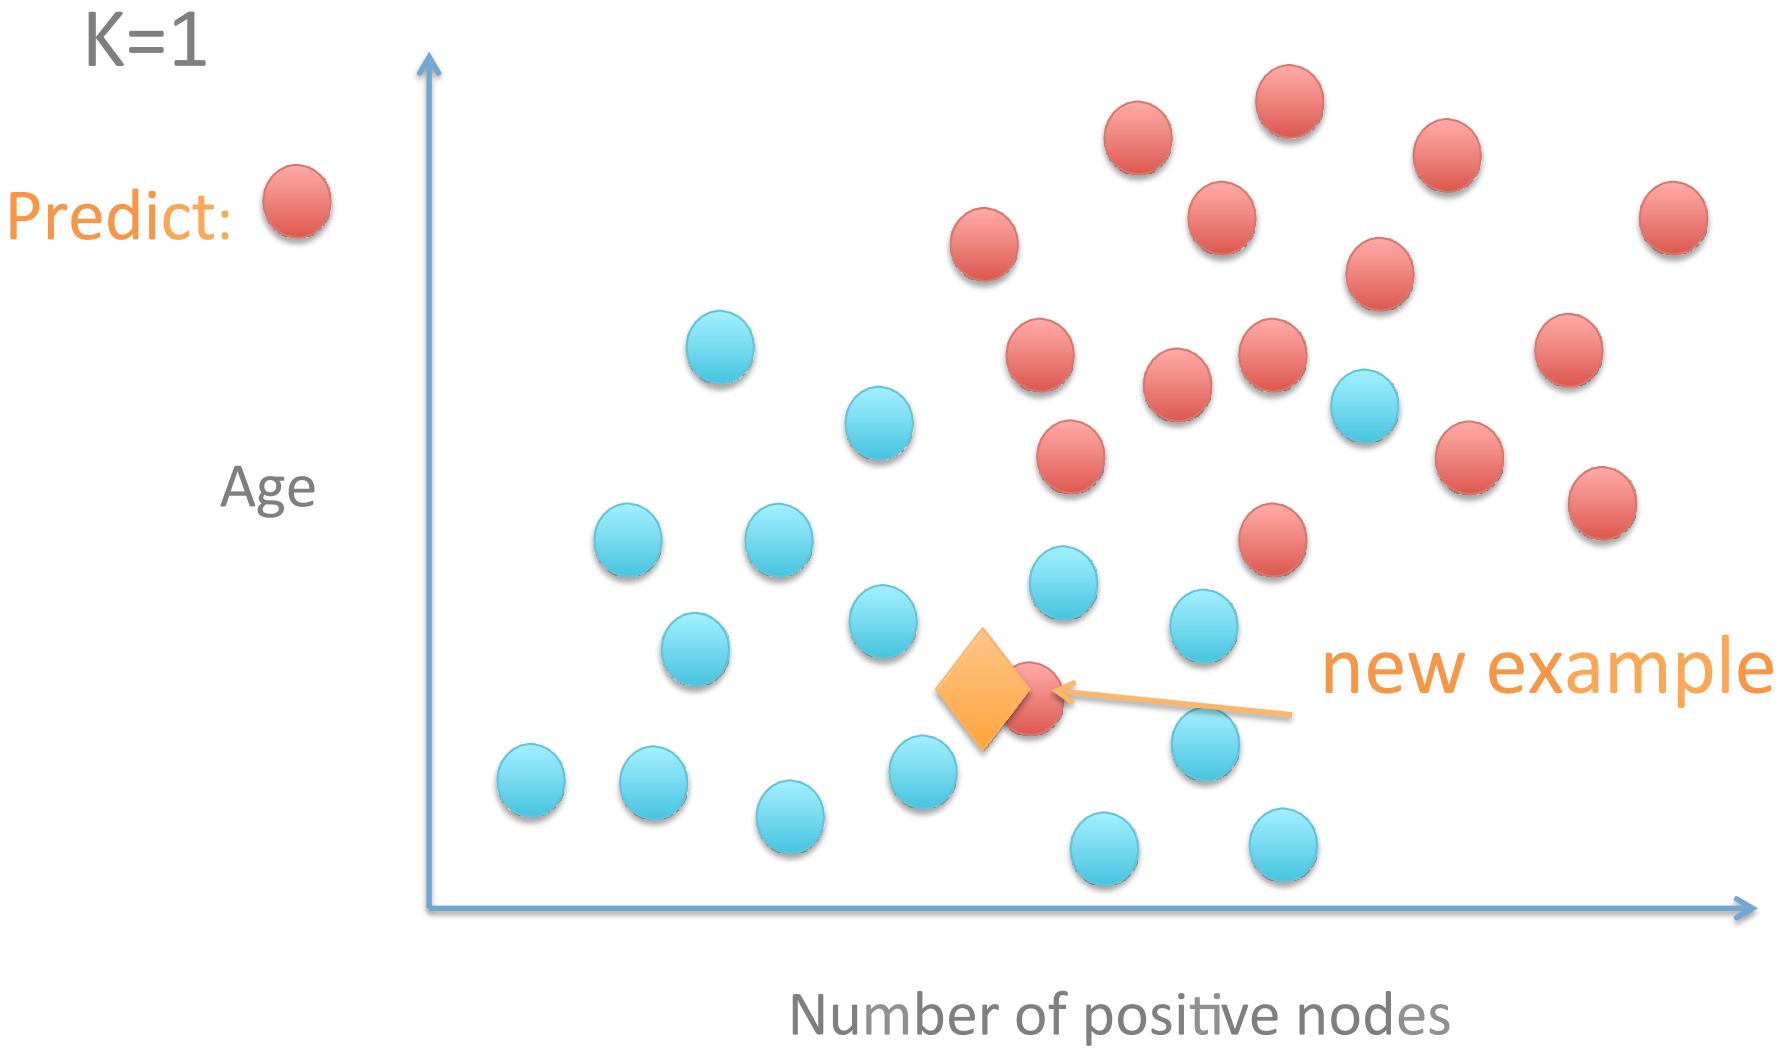

K Nearest Neighbors

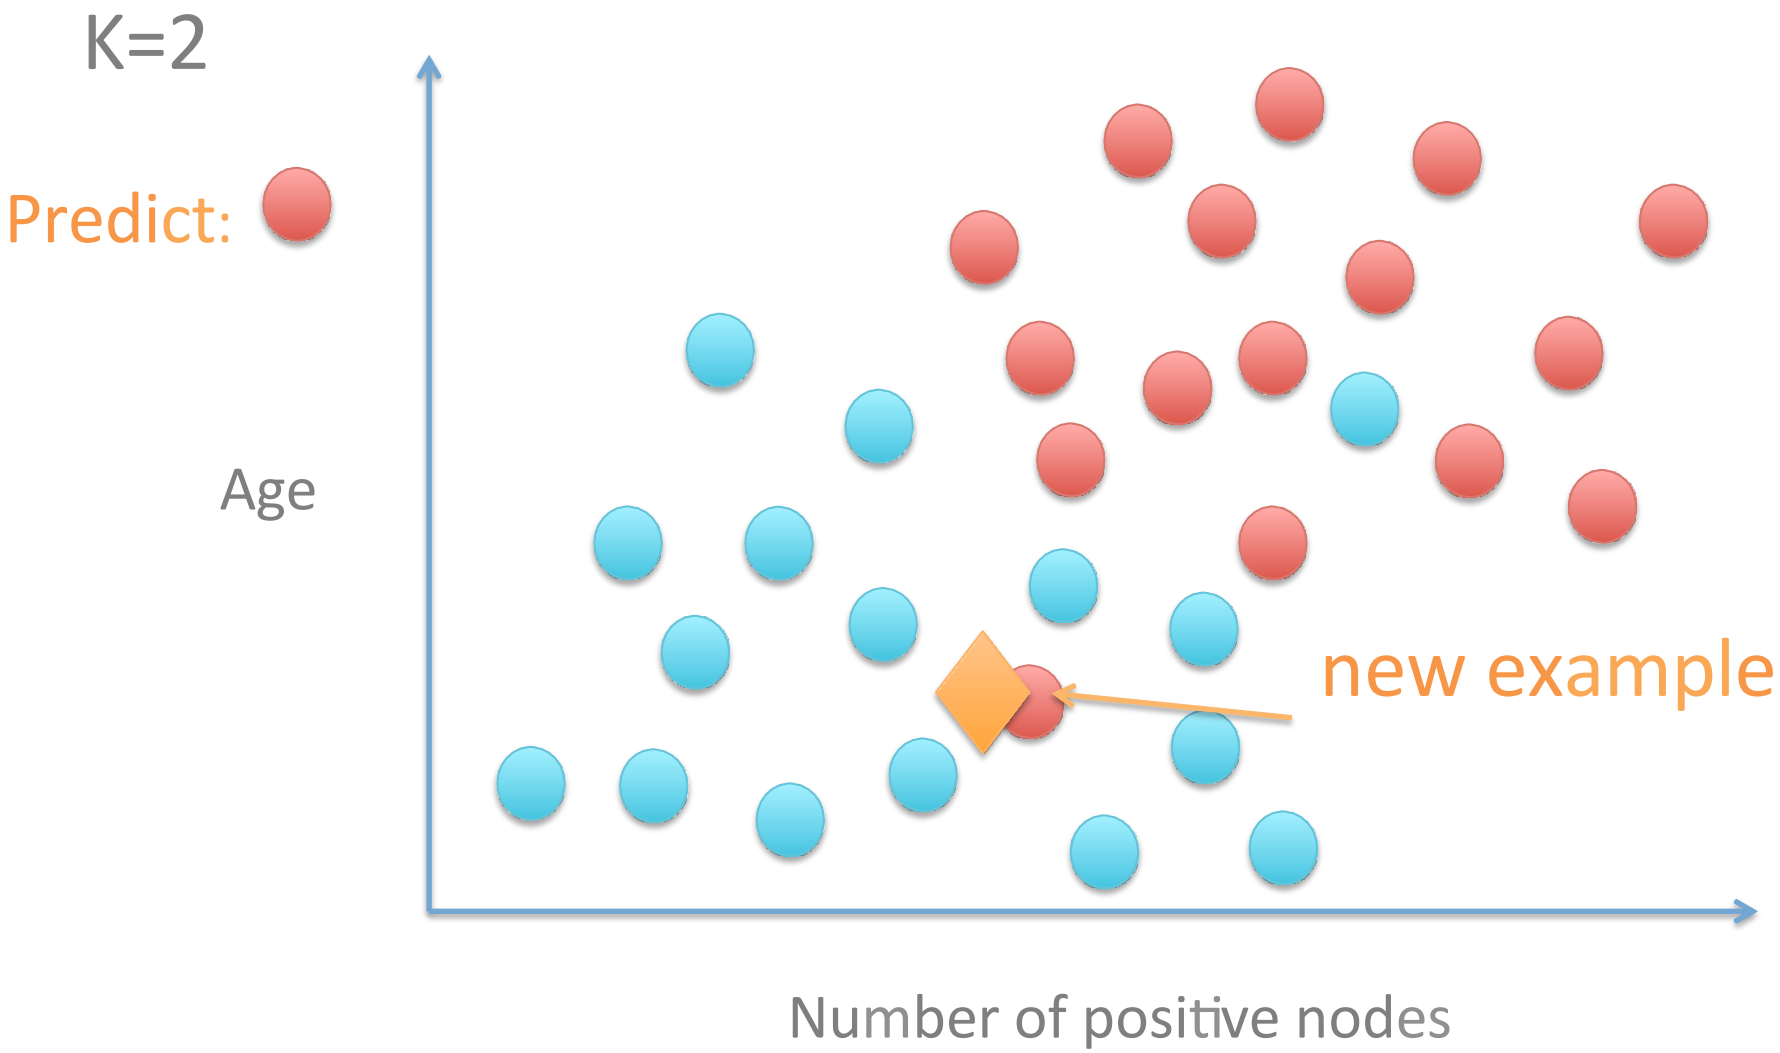

K Nearest Neighbors

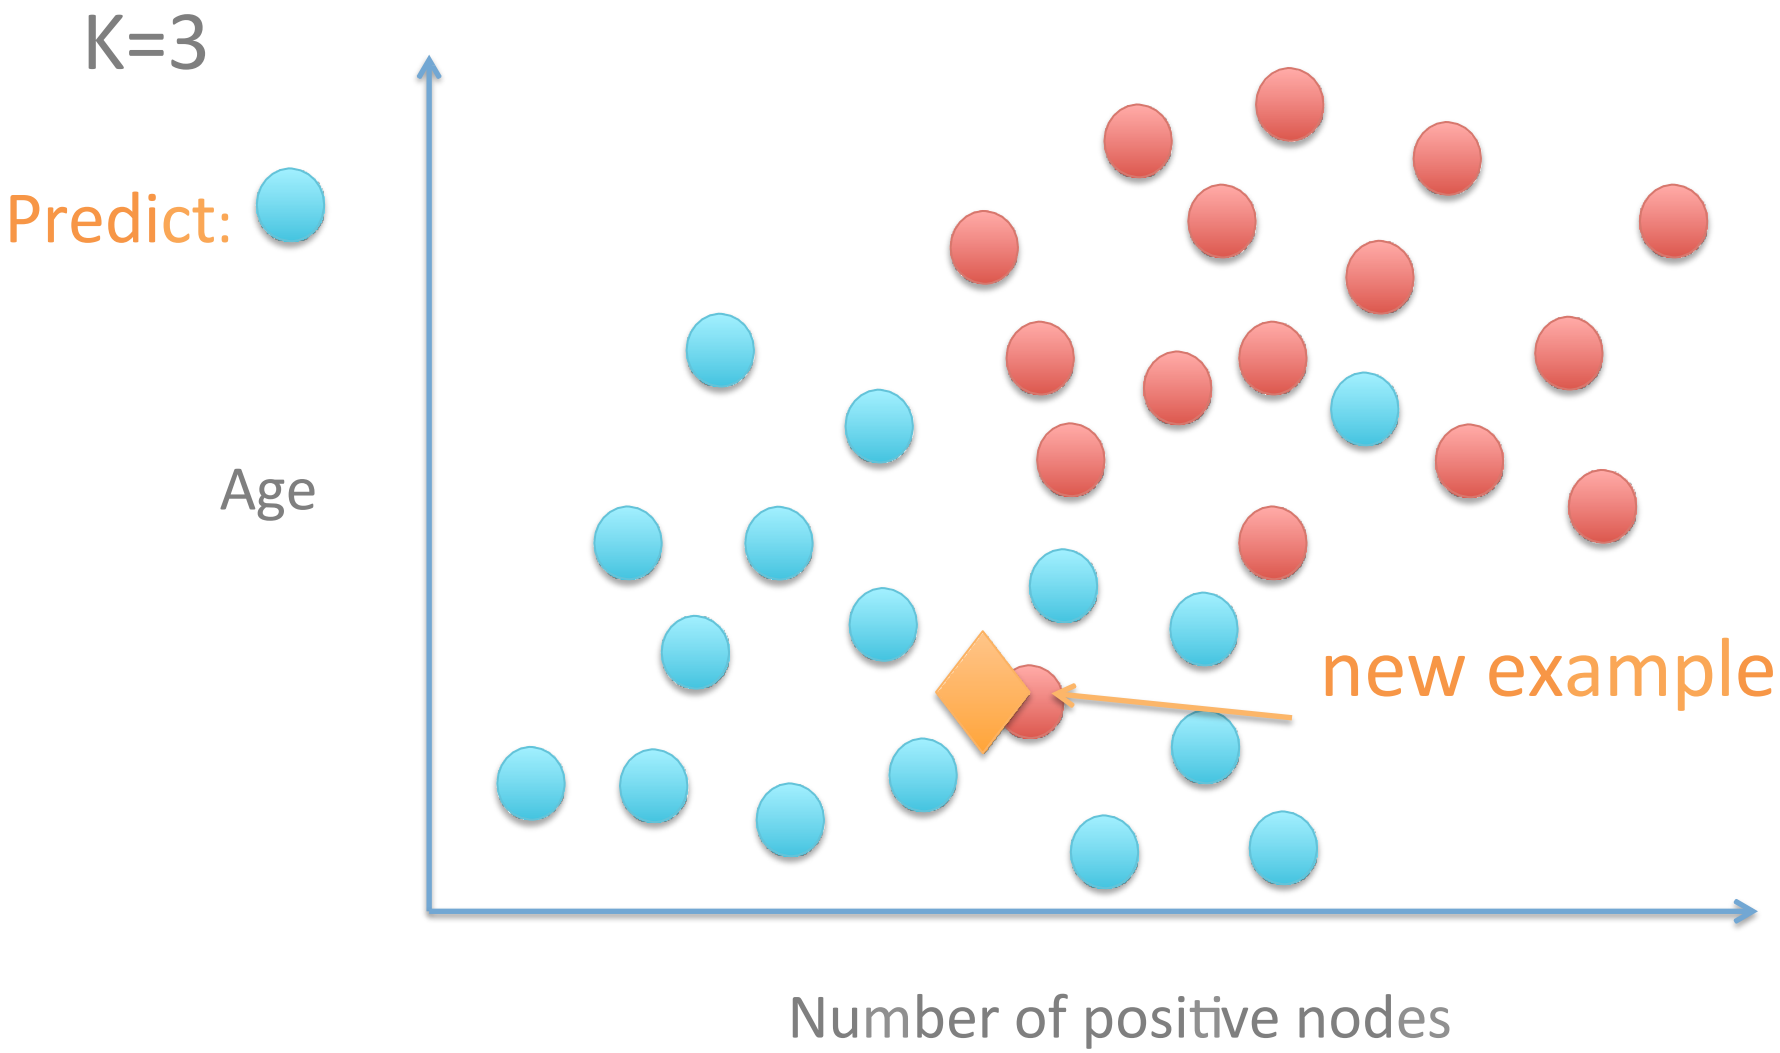

K Nearest Neighbors

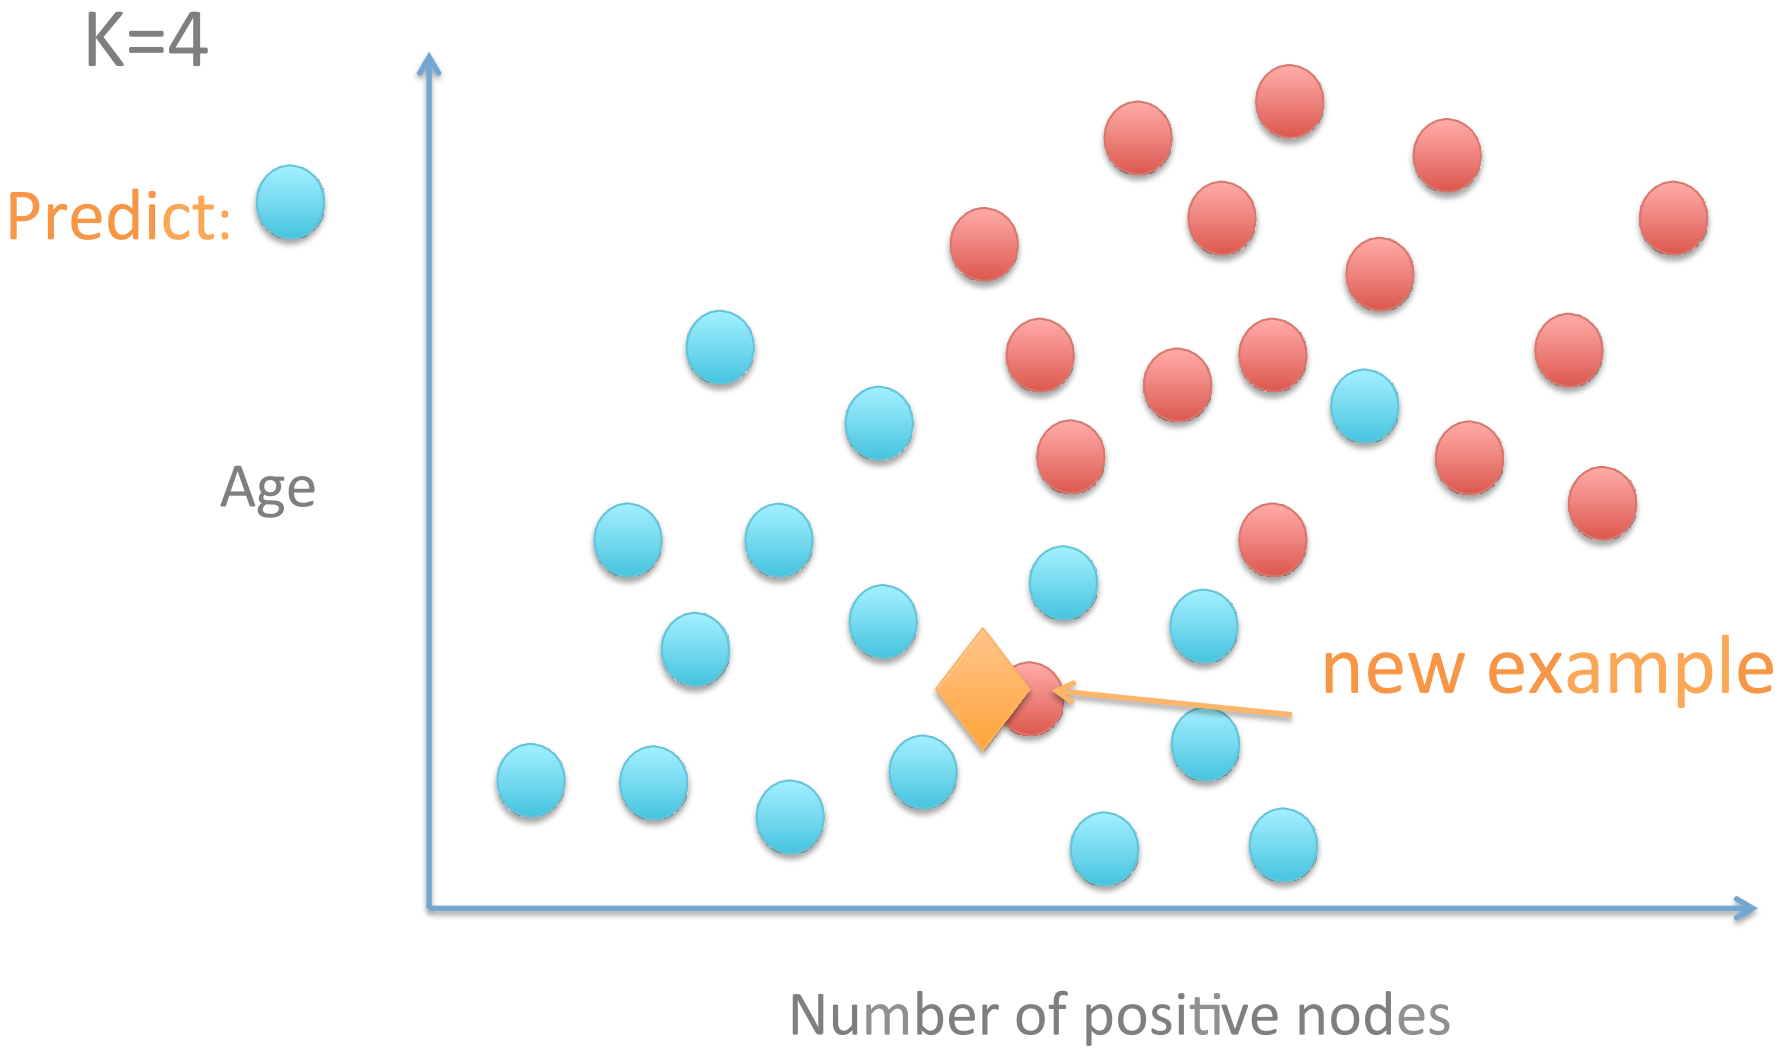

K Nearest Neighbors

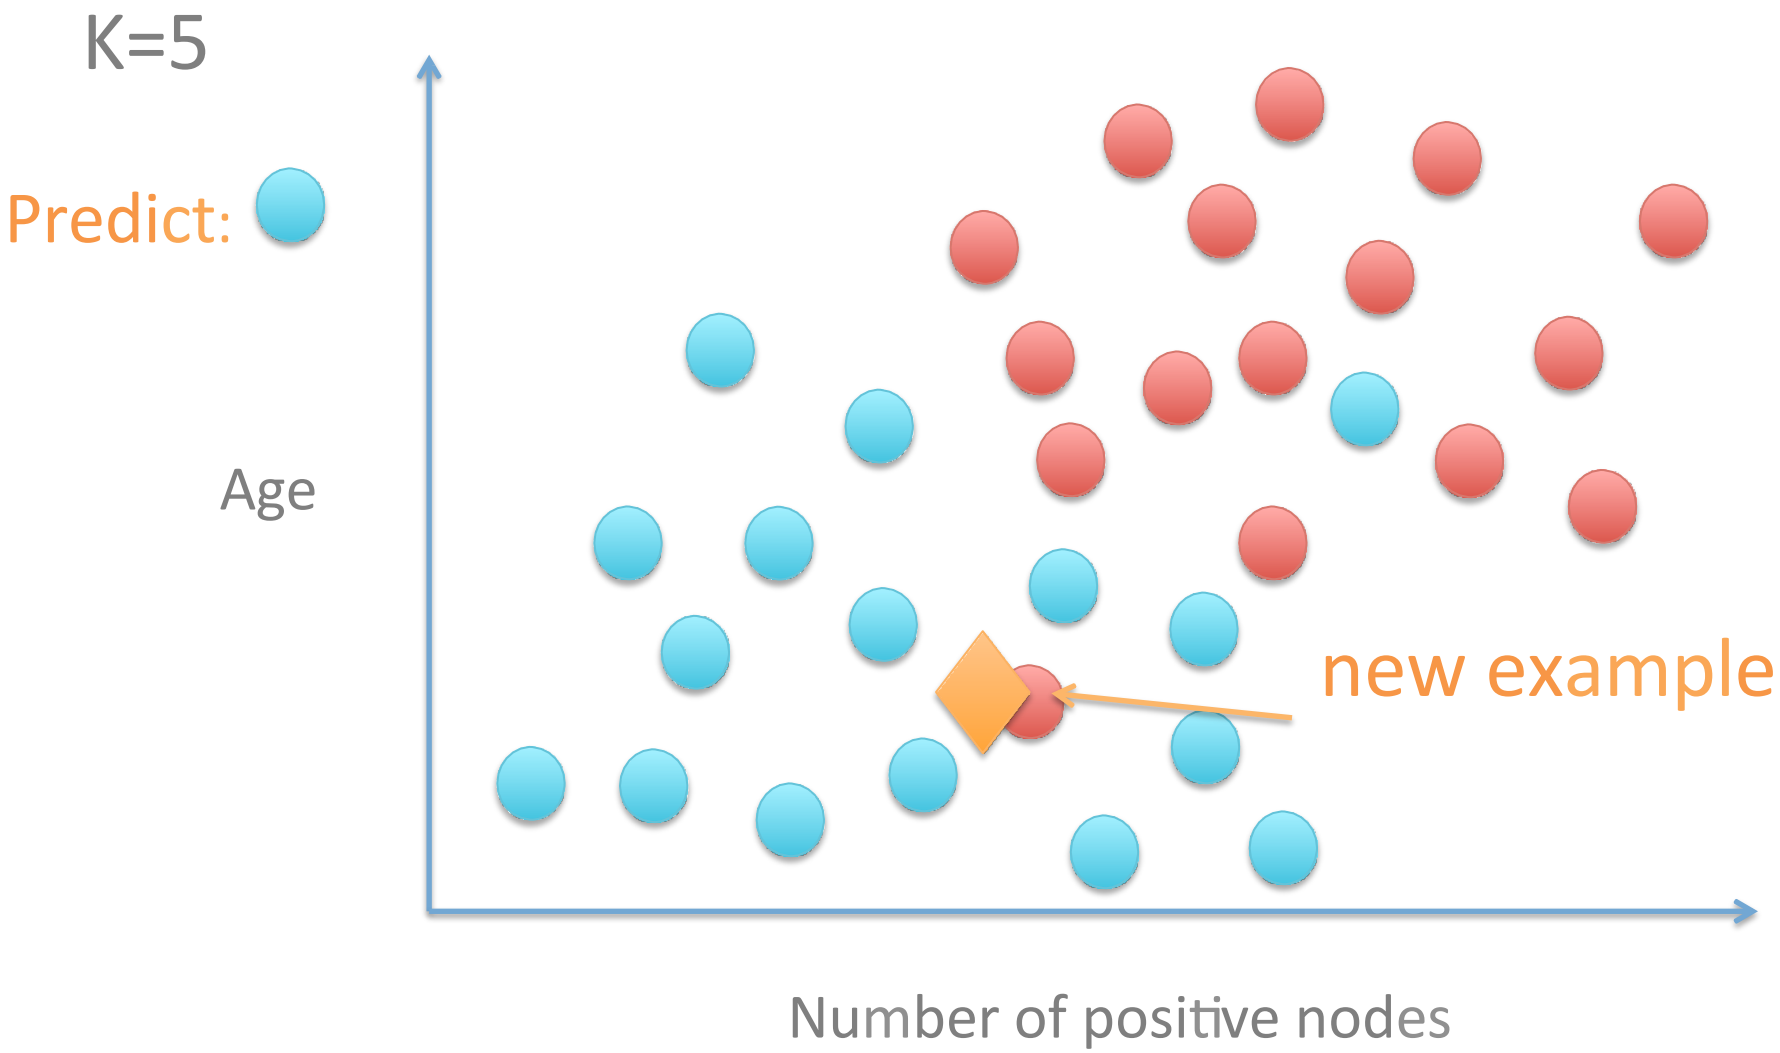

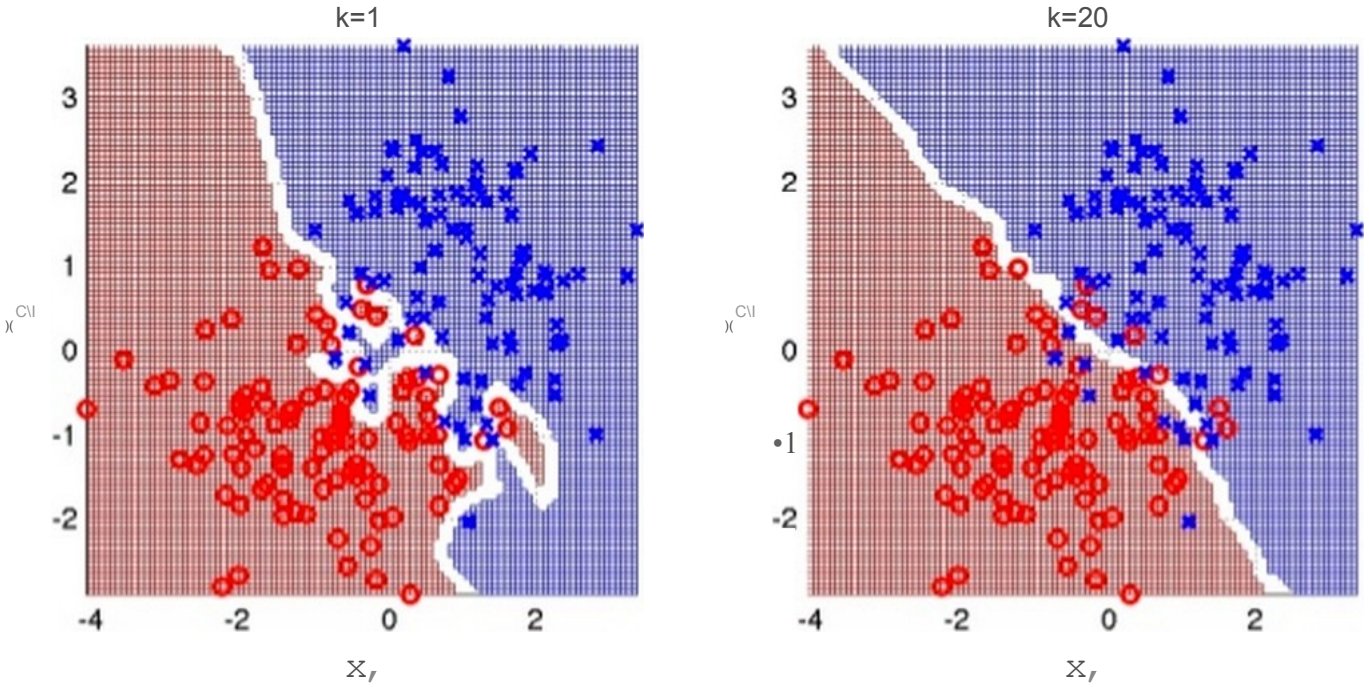
k=1
k=20
3
2
1
0
-1
-2
-4
-2
0
2

```
from sklearn.neighbors import KNeighborsClassifier
```